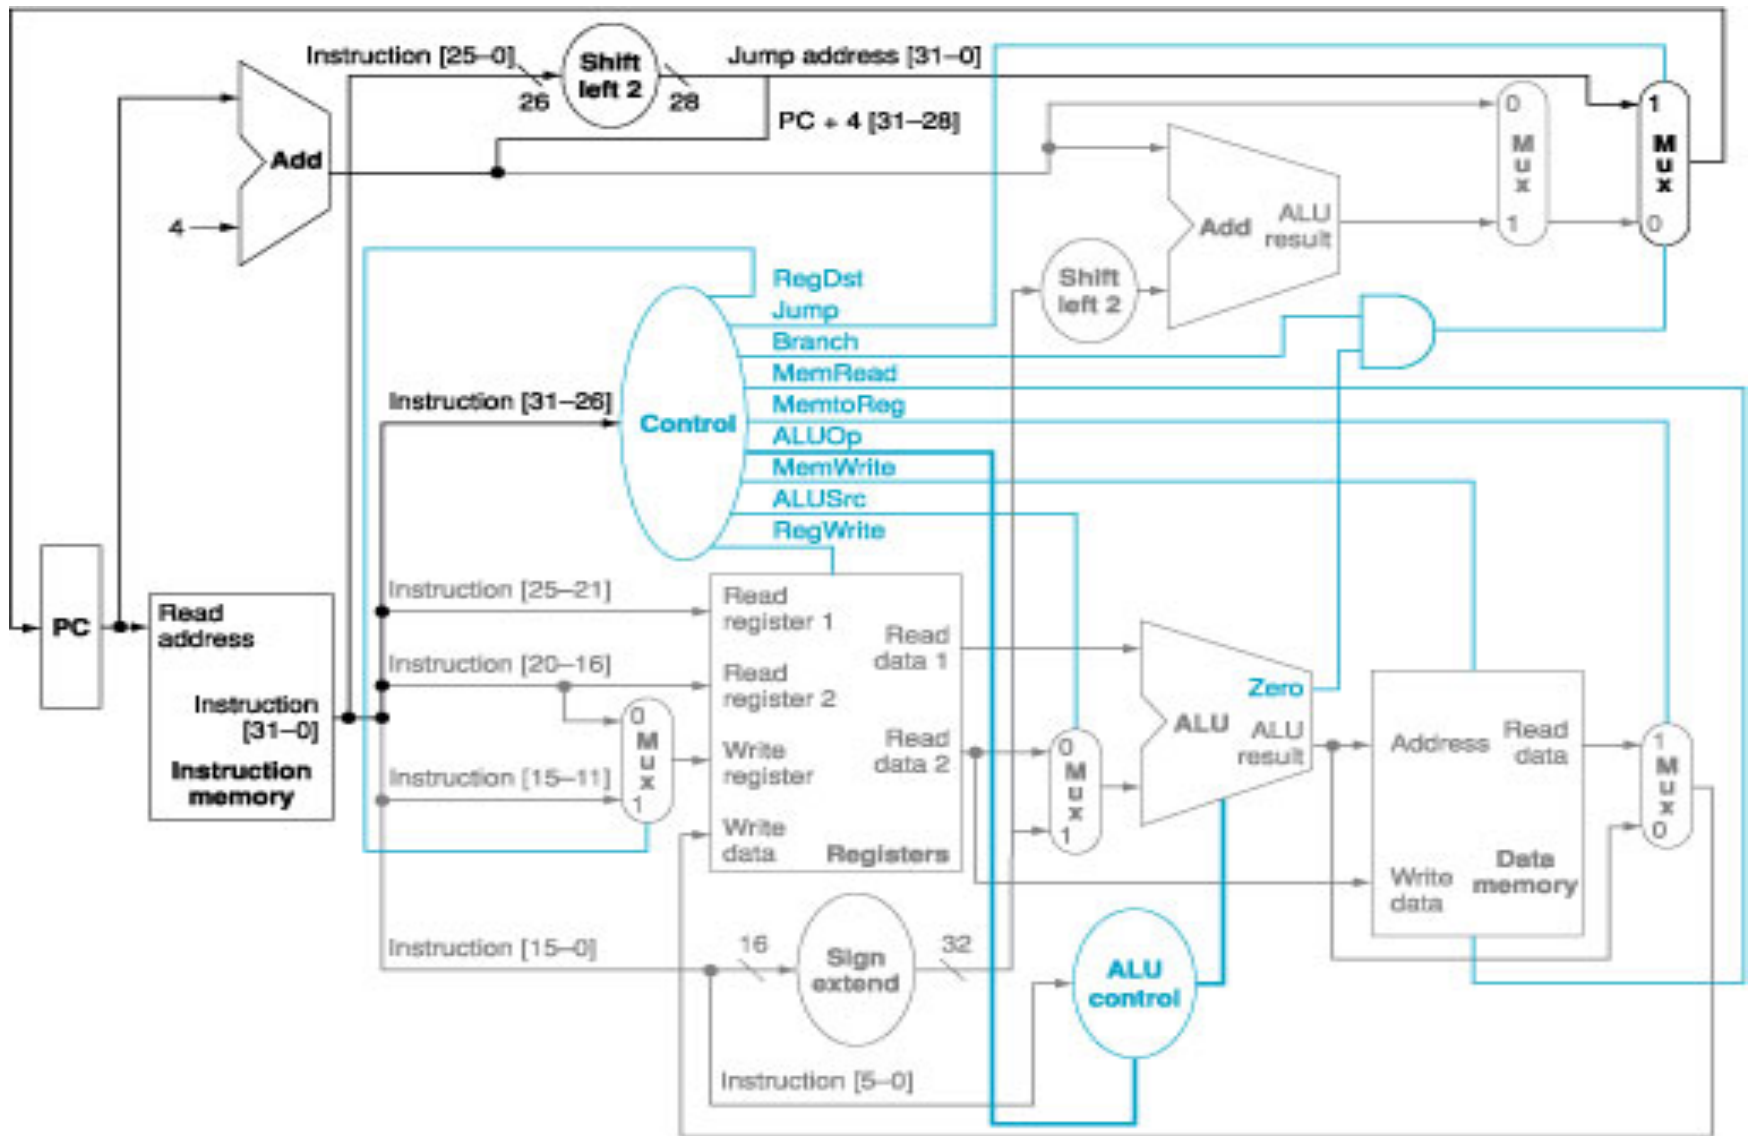


# Lecture 5

---

## Pipelining

# Single Cycle Implementations



# What's wrong with our CPI=1 processor?

Arithmetic & Logical



Load



← *Critical Path* →

Store

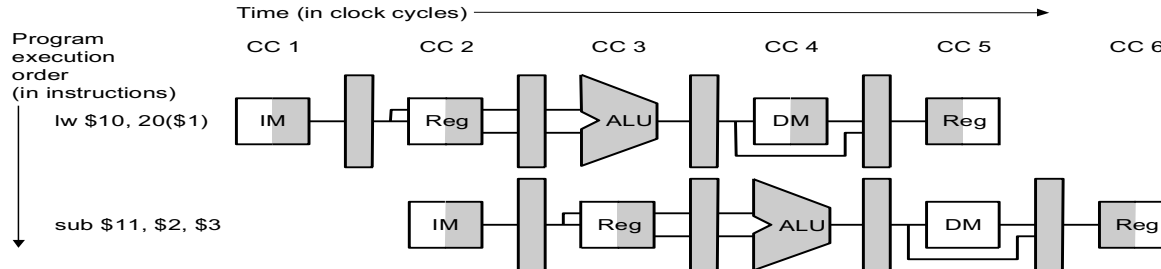


Branch



- Long Cycle Time
- All instructions take as much time as the slowest
- Real memory is not so nice as our idealized memory
  - cannot always get the job done in one (short) cycle

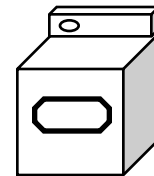
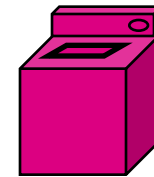

# Pipelining



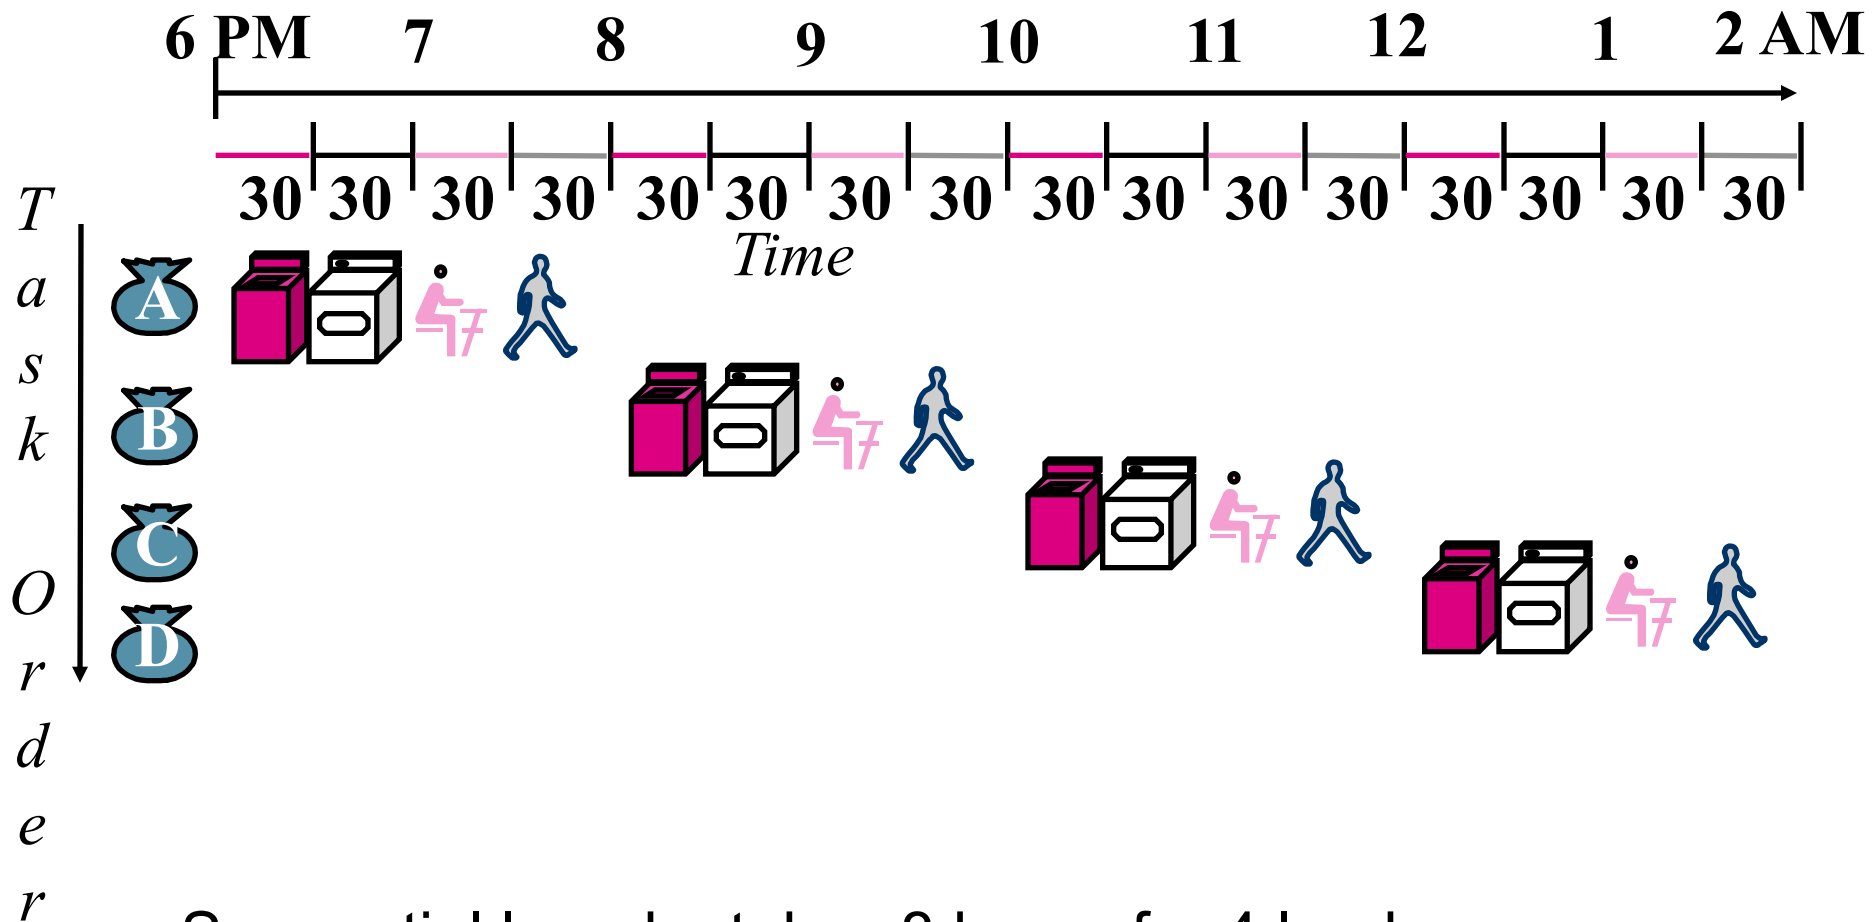
# Pipelining is Natural!

---

- Laundry Example
- Ann, Brian, Cathy, Dave each have one load of clothes to wash, dry, and fold
- Washer takes 30 minutes
- Dryer takes 30 minutes
- “Folder” takes 30 minutes
- Putting clothes into drawers takes 30 minutes

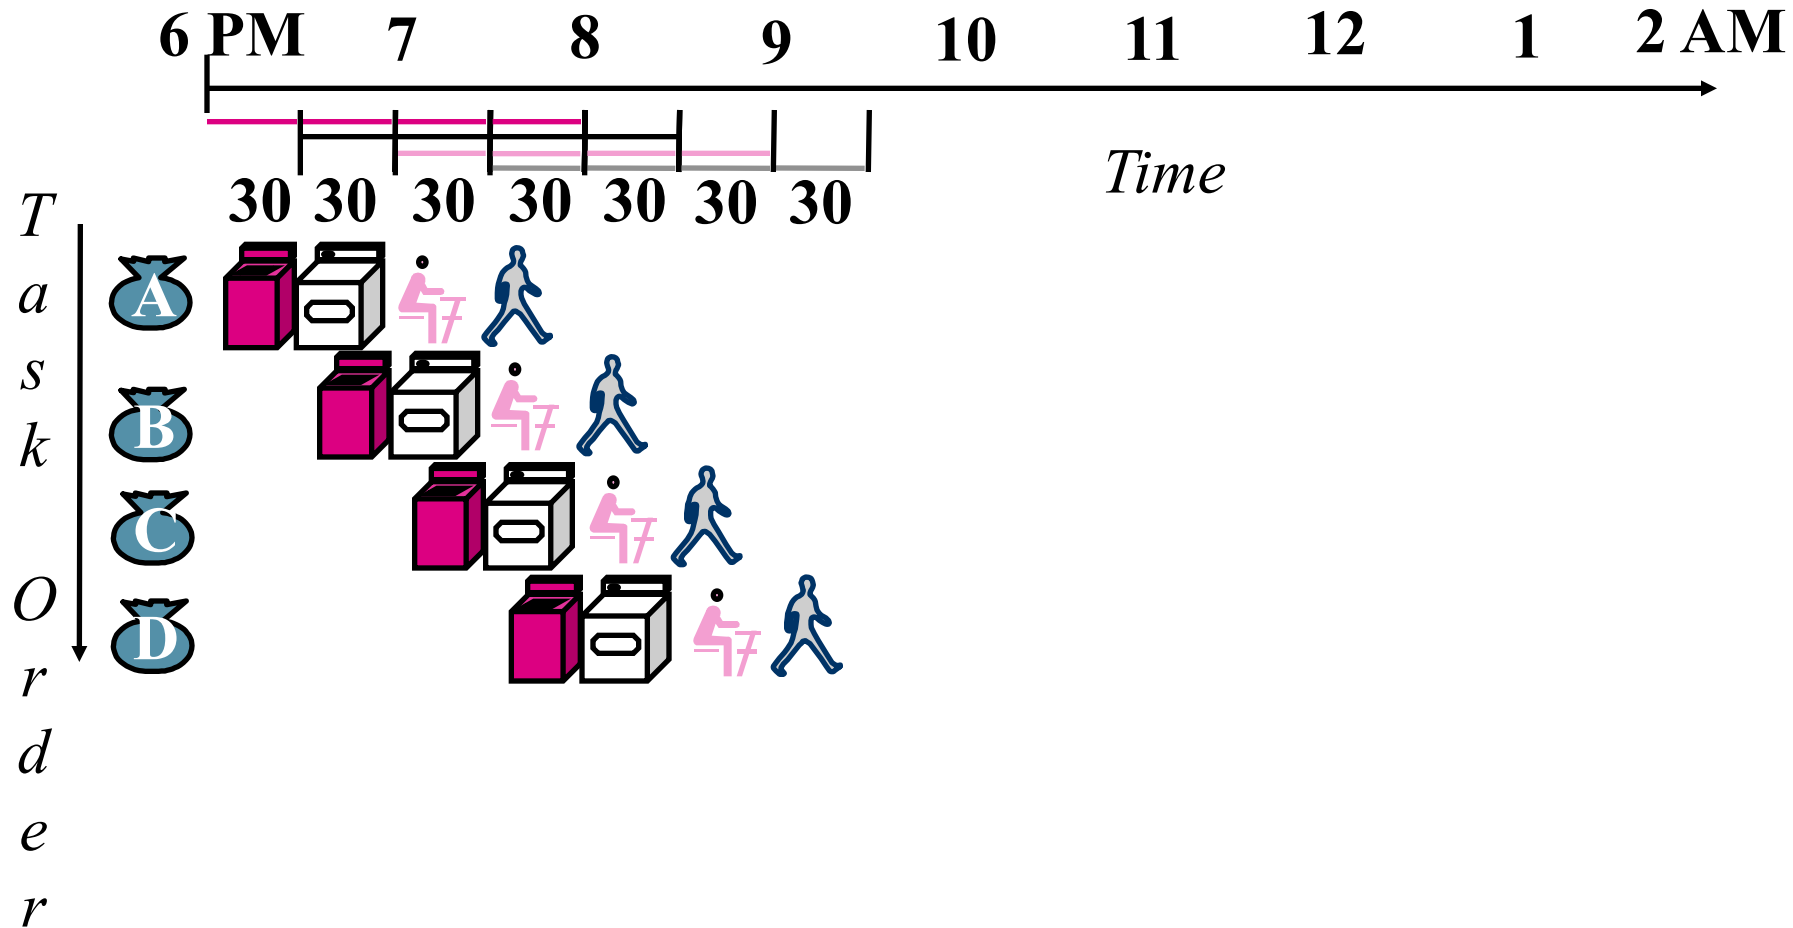


# Sequential Laundry




- Sequential laundry takes 8 hours for 4 loads
- If they learned pipelining, how long would laundry take?

# Pipelined Laundry: Start work ASAP



■ Pipelined laundry takes 3.5 hours for 4 loads!

# Pipelining Lessons



- Pipelining doesn't help **latency** of single task, it helps **throughput** of entire workload
- **Multiple** tasks operating simultaneously using different resources
- Pipeline rate limited by **slowest** pipeline stage
- Unbalanced lengths of pipe stages reduces speedup



# Pipelining in CPU

---


- Break the instruction into smaller steps
- Execute each step (instead of the entire instruction) in 1 clock cycle
  - Cycle time: time it takes to execute the longest step
  - Try to make all the steps have similar length

5 steps in MIPS

- Instruction Fetch
- Instruction Decode and Register Fetch
- Execution, Memory Address Computation, or Branch Completion
- Memory Access or R-type instruction completion
- Write-back

# The Five Stages of Load

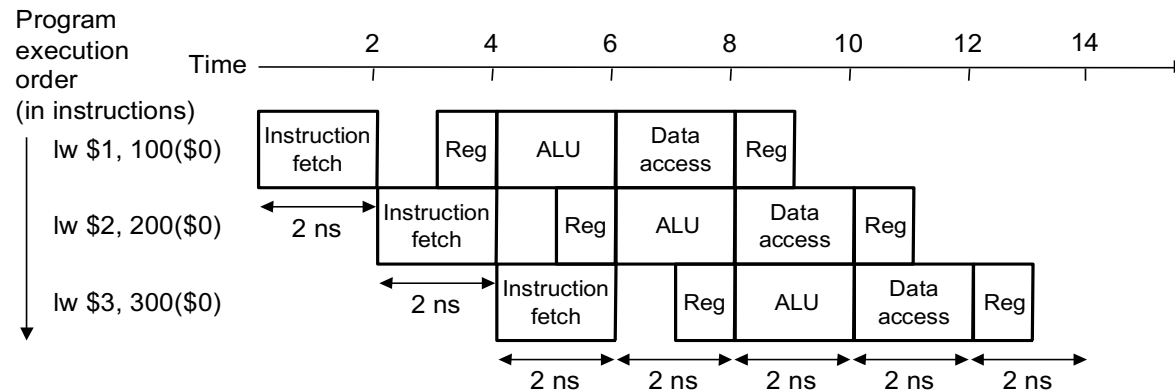
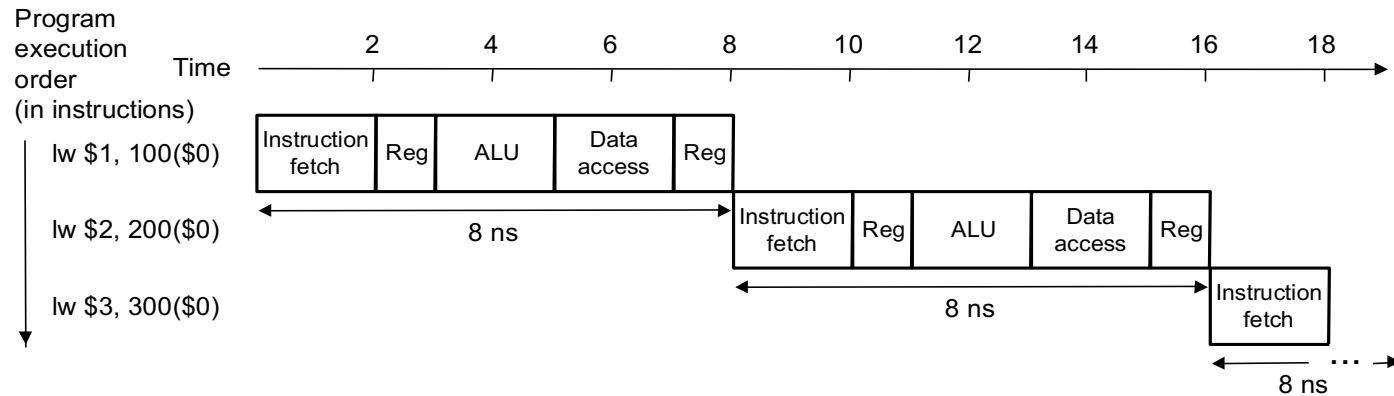
---



- **Ifetch: Instruction Fetch**
  - Fetch the instruction from the Instruction Memory
- **Reg/Dec: Registers Fetch and Instruction Decode**
- **Exec: Calculate the memory address**
- **Mem: Read the data from the Data Memory**
- **Wr: Write the data back to the register file**

# Pipelining

## ■ Improve performance by increasing instruction throughput




Time to execute 4 load instructions:

Single cycle:  $8 \times 4 = 32$  ns


Pipeline:  $4 \times 1 + 4$  (time to drain the pipeline) = 8 Cycle = 16 ns

# Single Cycle vs. Pipeline

## Single Cycle Implementation:



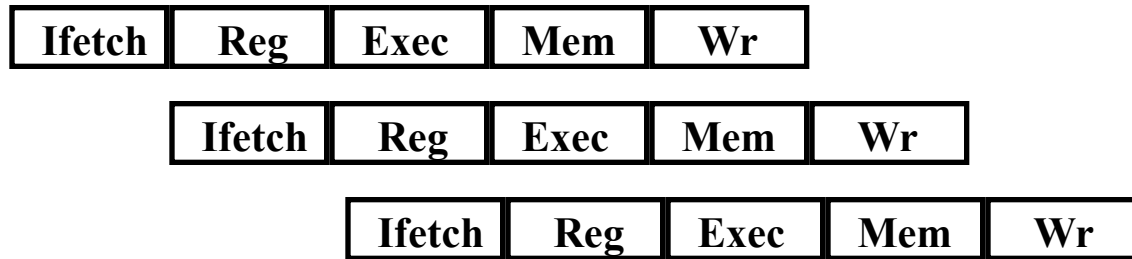
## Pipeline Implementation:



# Ideal Pipeline Performance

---

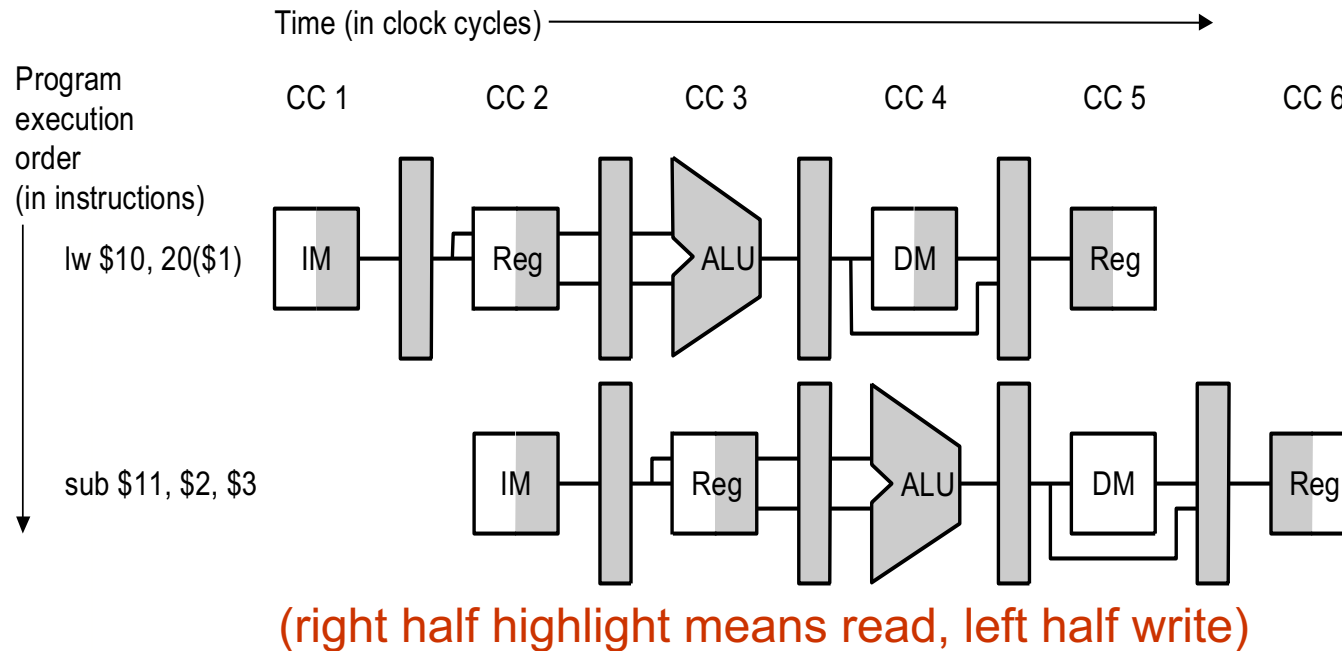
- Ideal CPI = 1
- Time to execute n instruction
  - $1 \times n + \text{time to drain the pipeline}$



- Ideal speedup from pipelining == # of pipeline stages
  - If the stages are perfectly balanced, and a large number of instructions

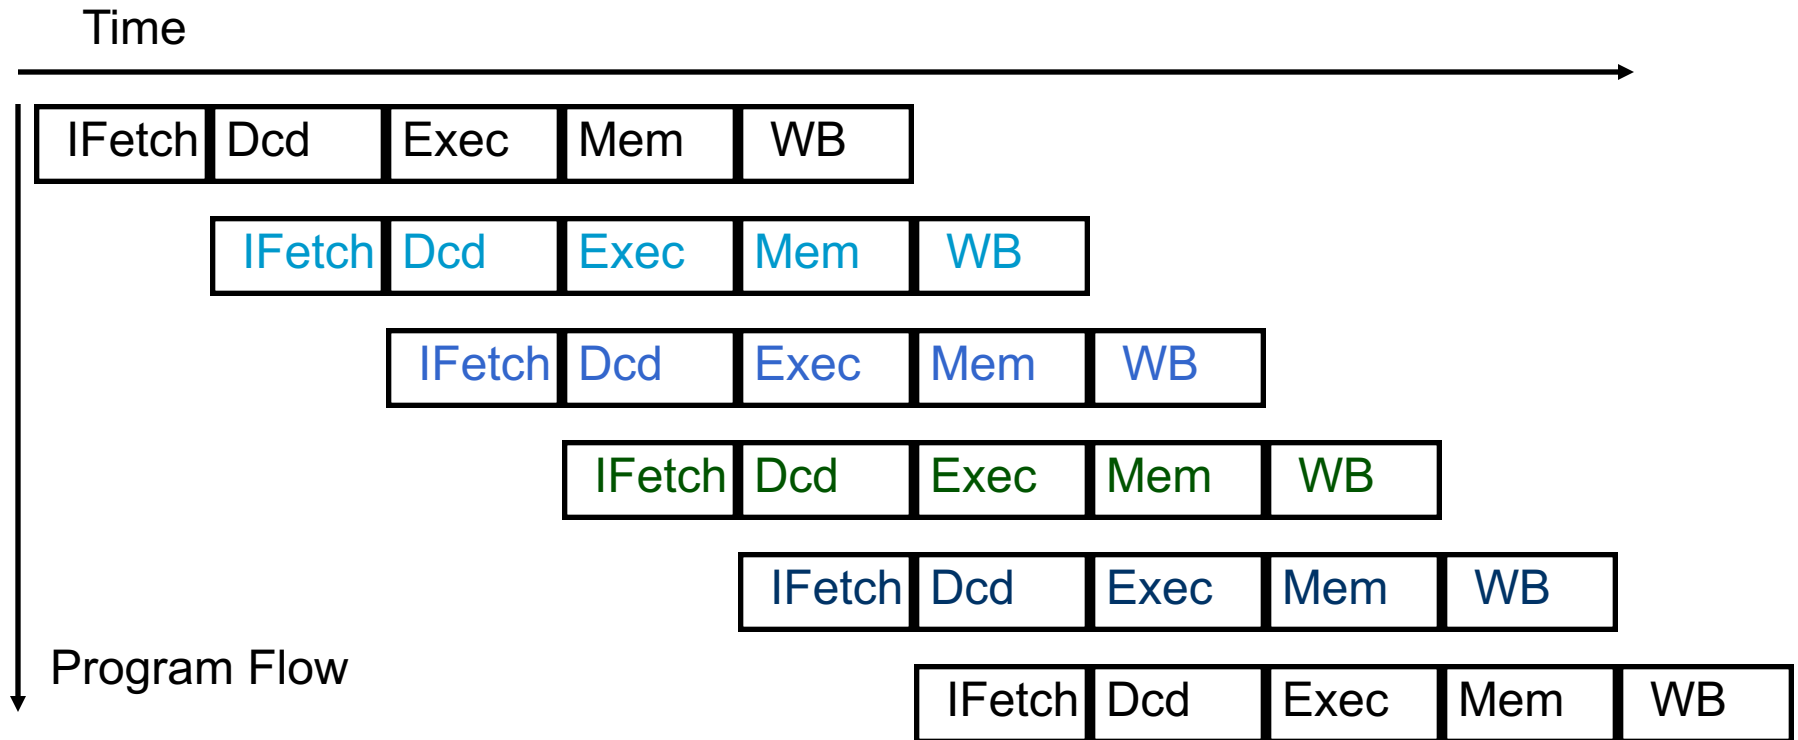
$$\text{Time between instructions}_{\text{pipelined}} = \frac{\text{Time between instructions}_{\text{nonpipelined}}}{\text{Number of pipe stages}}$$

# Graphically Representing Pipelines

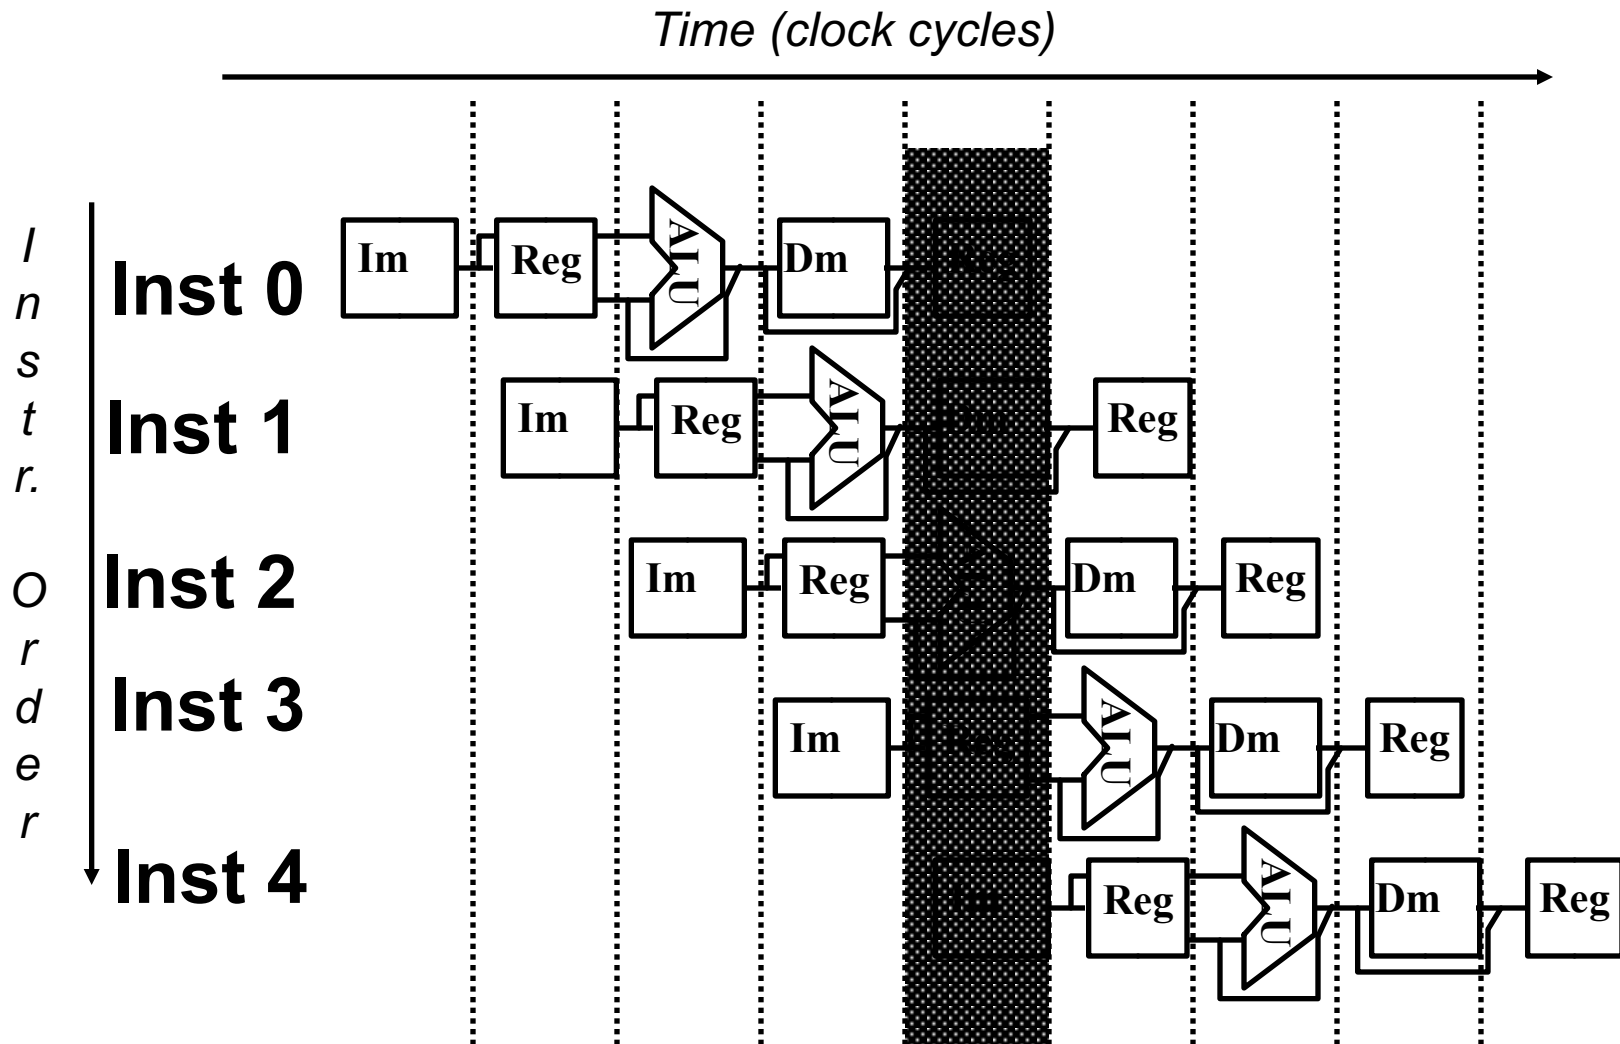


- Can help with answering questions like:
  - how many cycles does it take to execute this code?
  - what is the ALU doing during cycle 4?
  - use this representation to help understand datapaths

# Conventional Pipelined Execution Representation



# Instruction Pipelining





# Can pipelining get us into trouble?

---

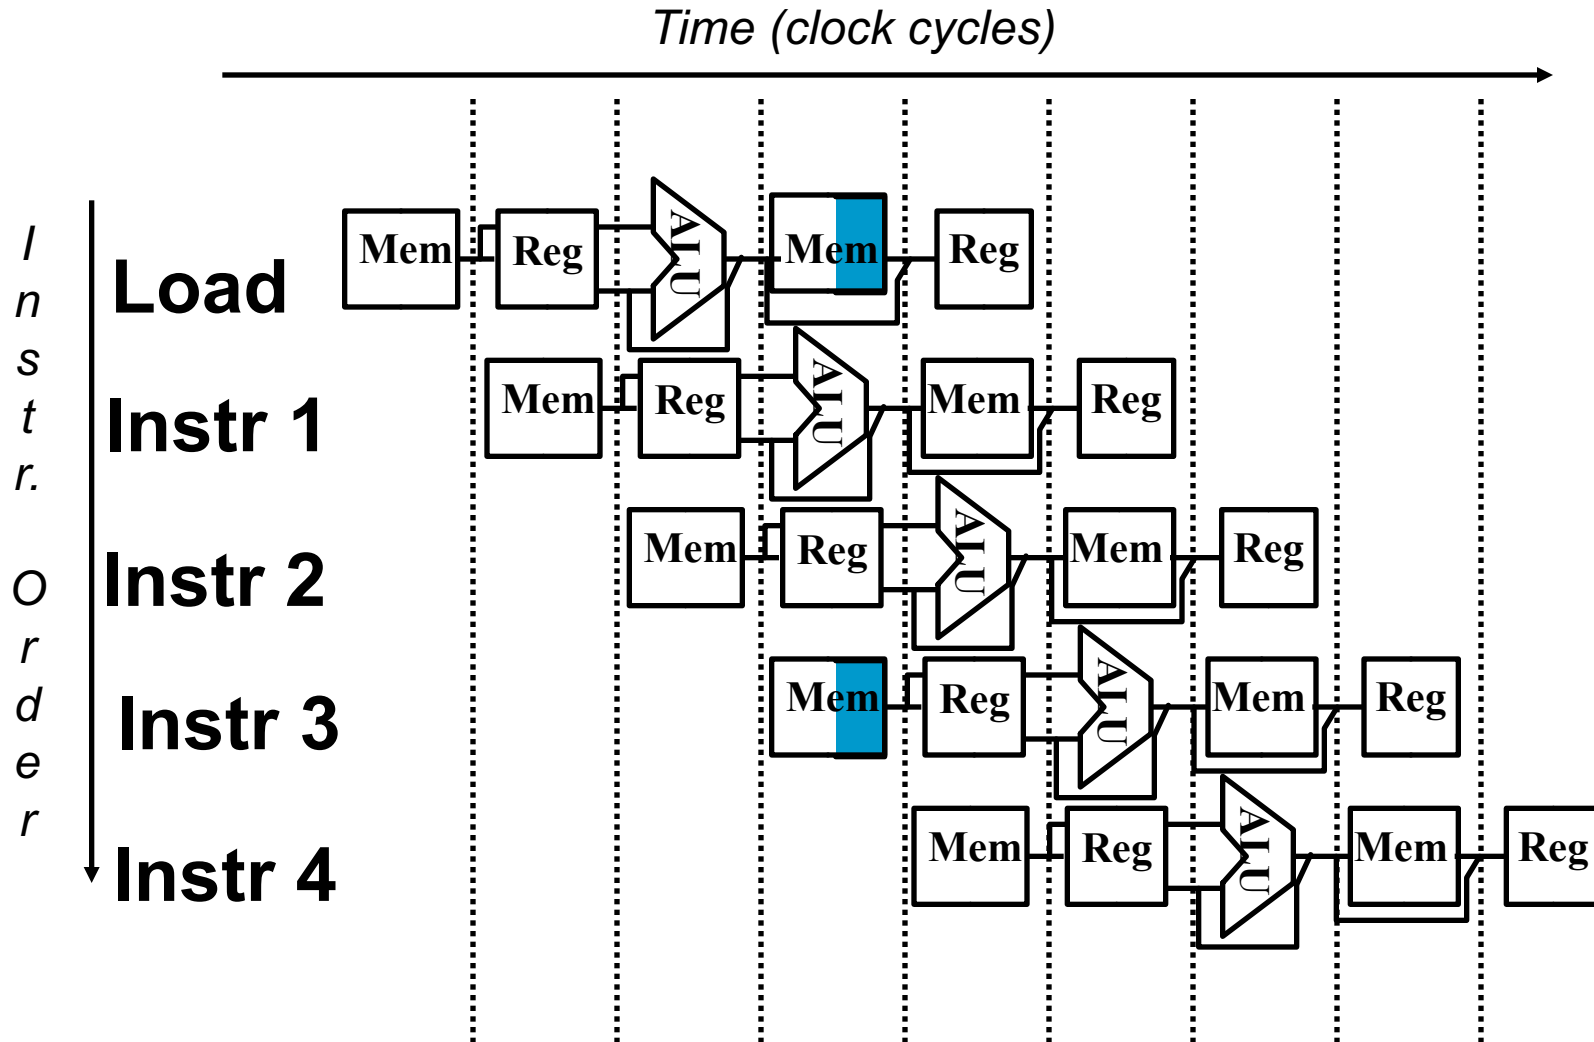
## ■ Yes: Pipeline Hazards

- **structural hazards**: attempt to use the same resource two different ways at the same time
  - E.g., combined washer/dryer would be a structural hazard or folder busy doing something else (watching TV)
- **data hazards**: attempt to use item before it is ready
  - E.g., one sock of pair in dryer and one in washer; can't fold until get sock from washer through dryer
  - instruction depends on result of prior instruction still in the pipeline
- **control hazards**: attempt to make a decision before condition is evaluated
  - E.g., washing football uniforms and need to get proper detergent level ; need to see after dryer before next load in
  - branch instructions

## ■ Can always resolve hazards by **waiting**

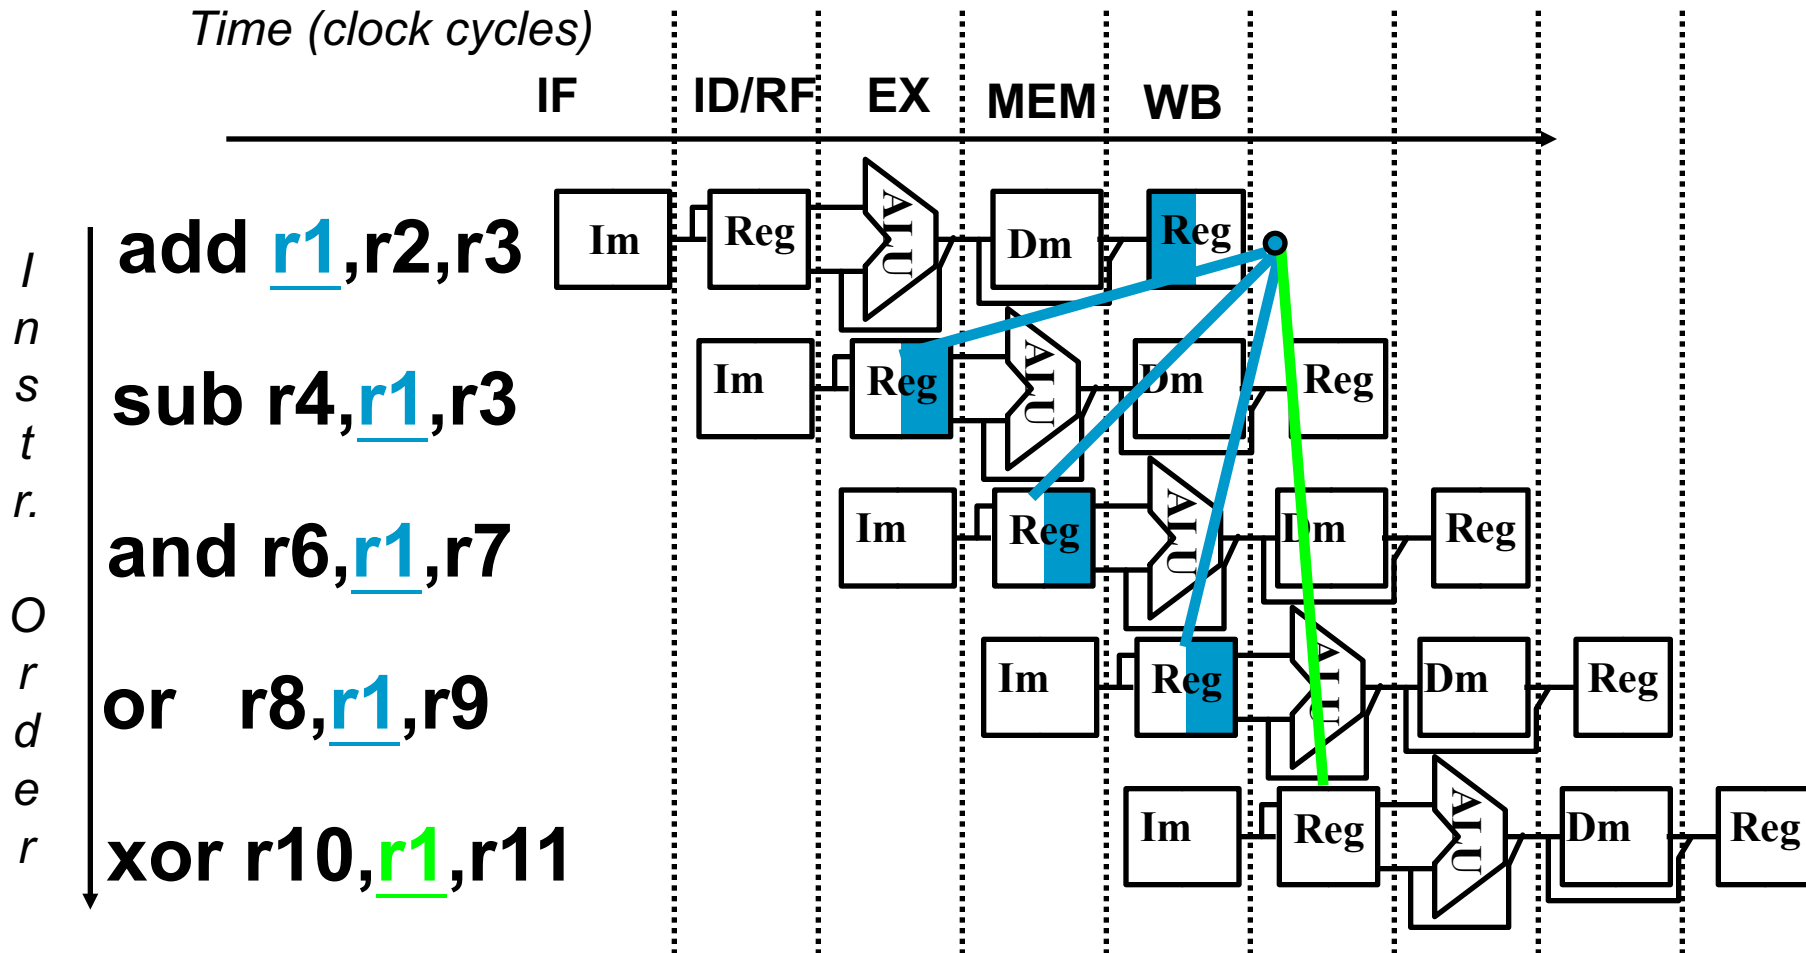
- pipeline control must detect the hazard
- take action (or delay action) to resolve hazards

# Single Memory is a Structural Hazard

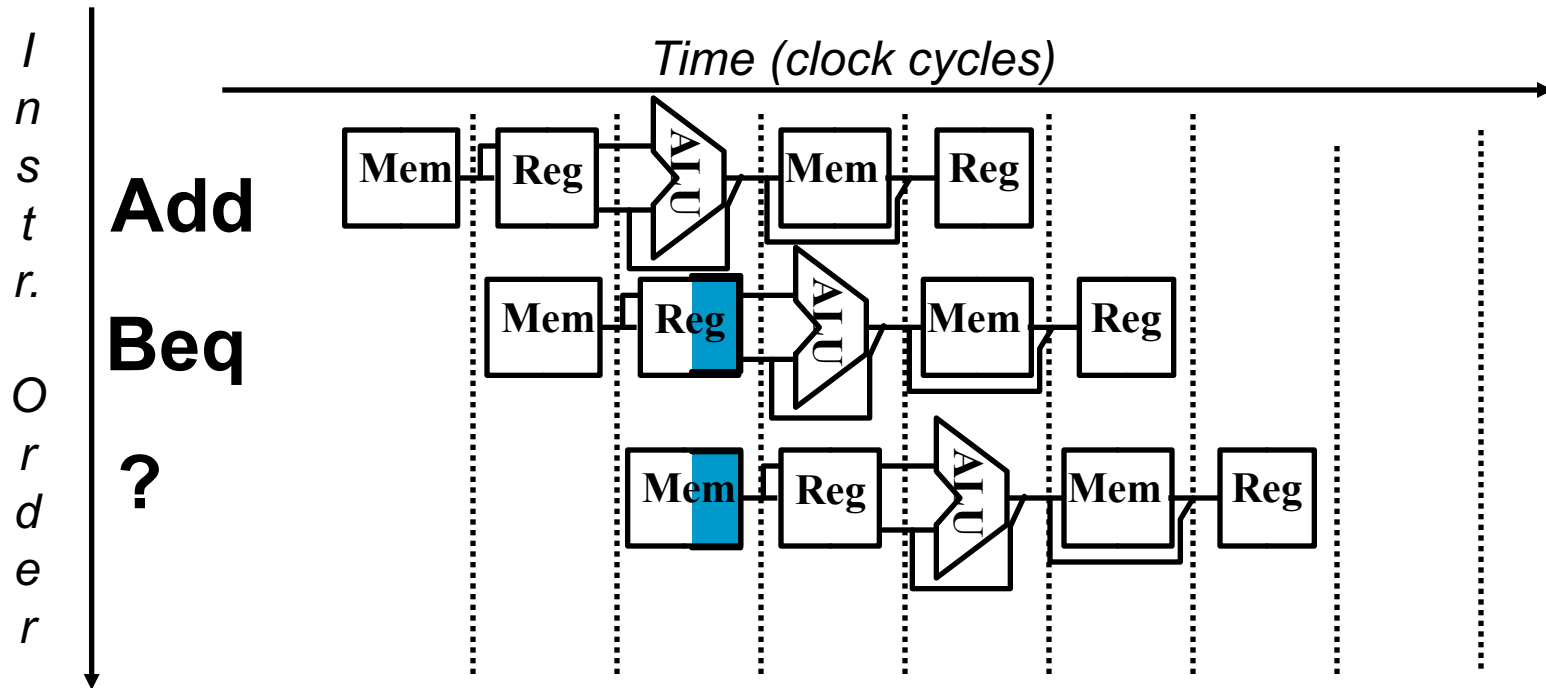


# Data Hazard on r1:

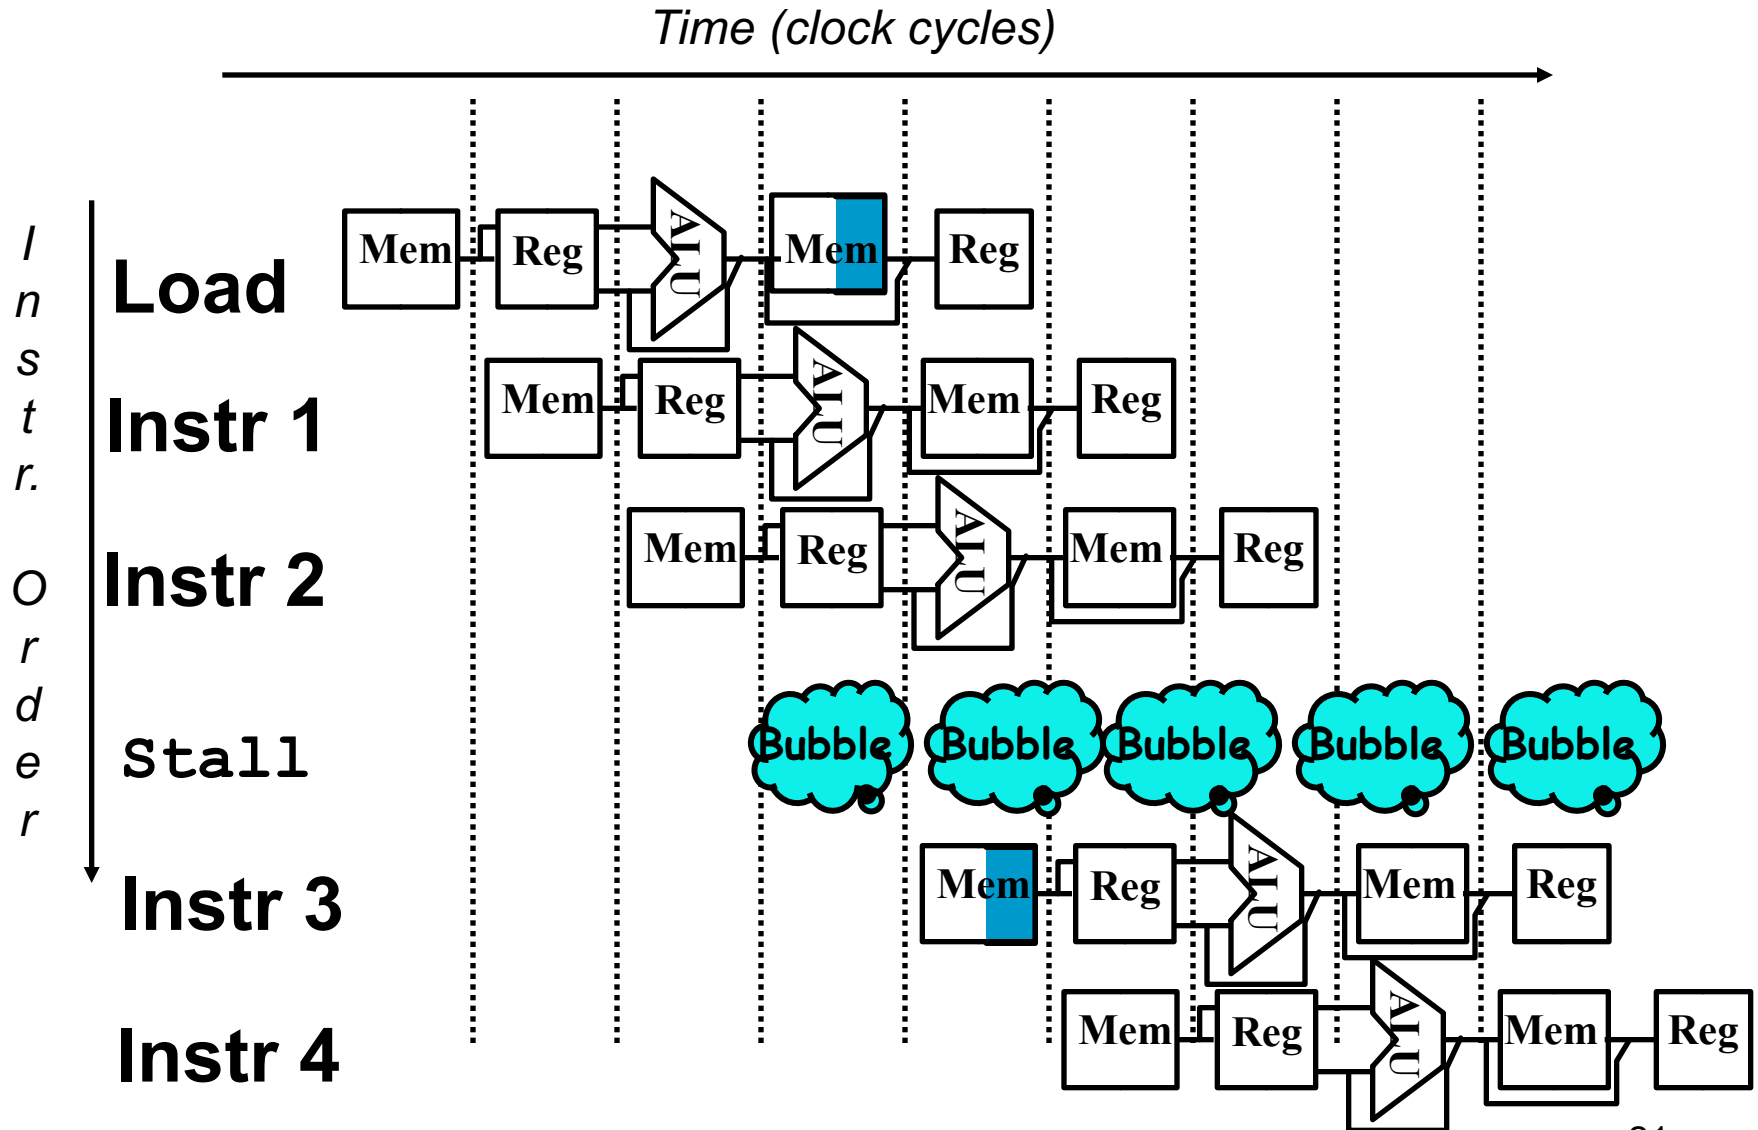
- **Dependencies** backwards in time are **hazards**



# Control Hazard

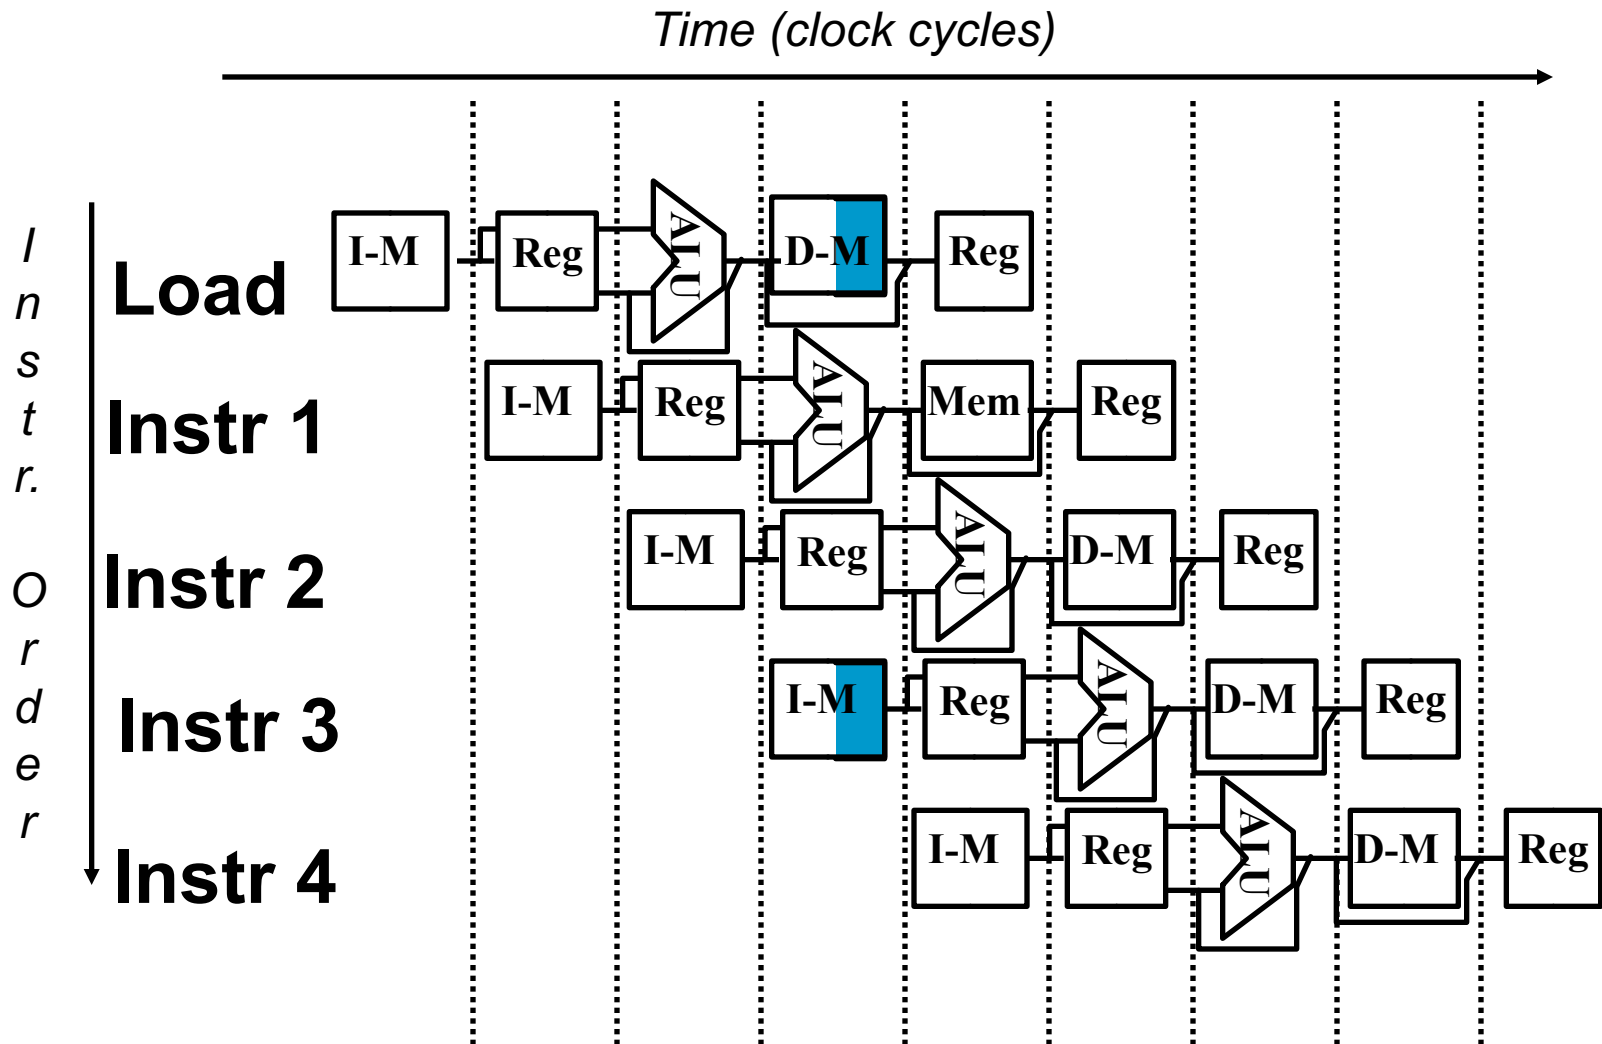


# How to solve structure hazards: (1) Stall




# How to solve structure hazards: (2)

## Split instruction and data memory



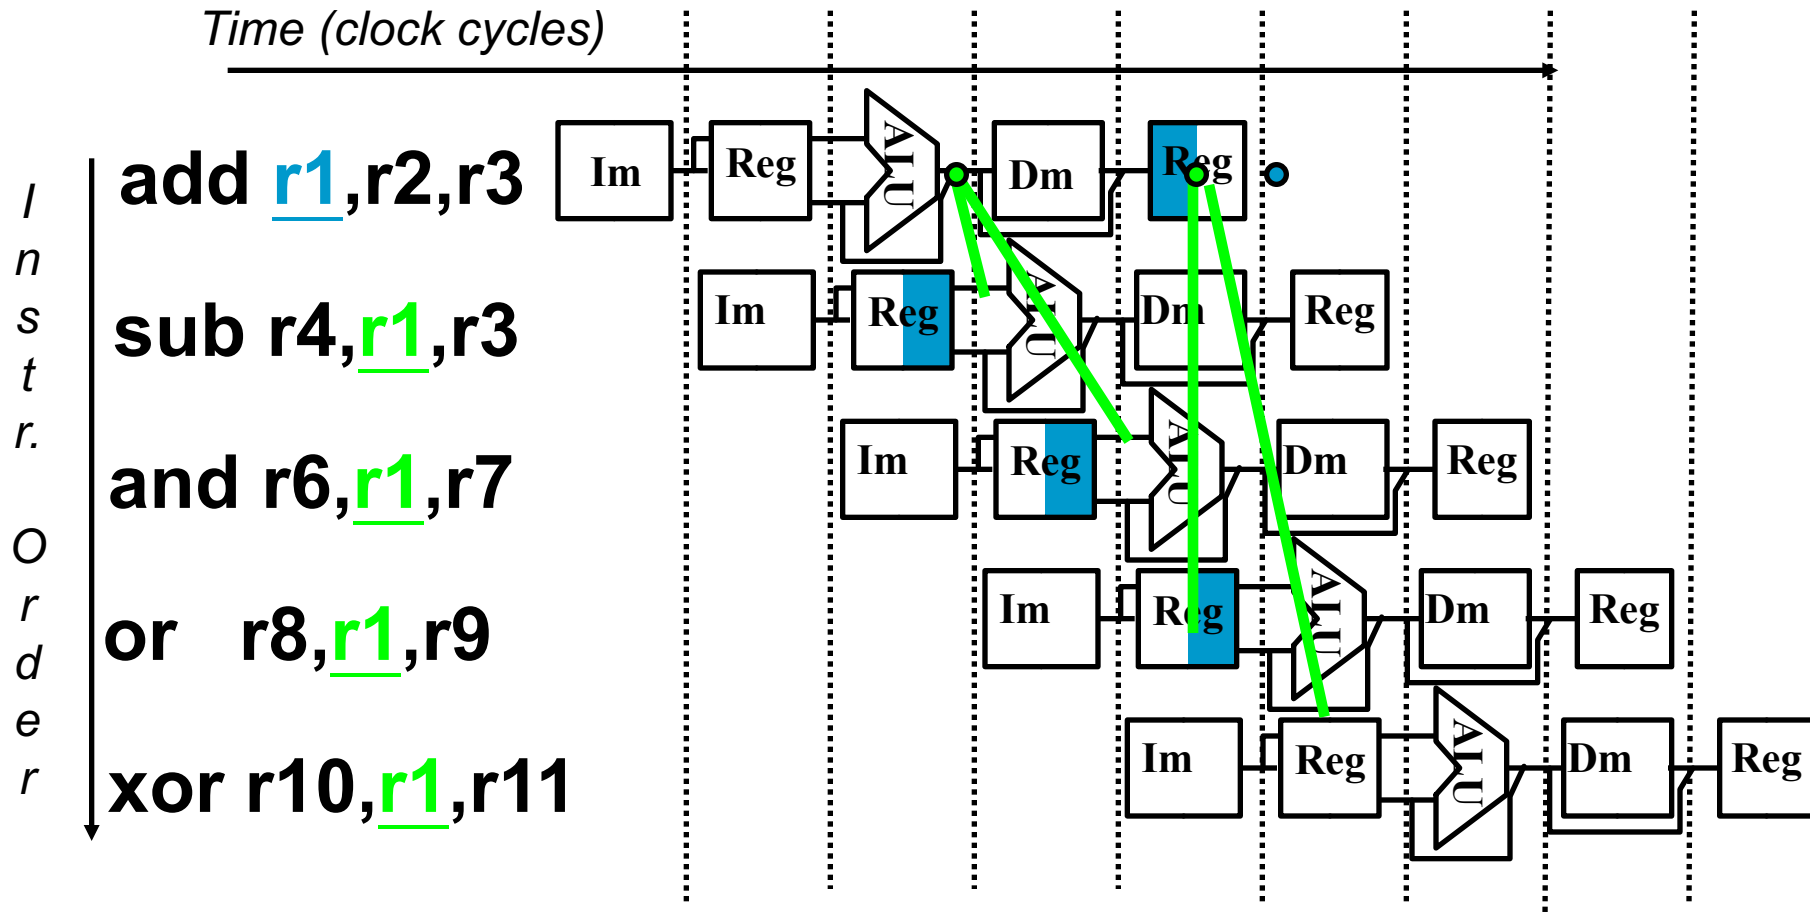
# Data Hazard on r1:

- Dependencies backwards in time are hazards



# Data Hazard Solution:

- “Forward” result from one stage to another




- “or” OK if register writes occur in the first half of the clock cycle, and register reads in the last half of the clock cycle

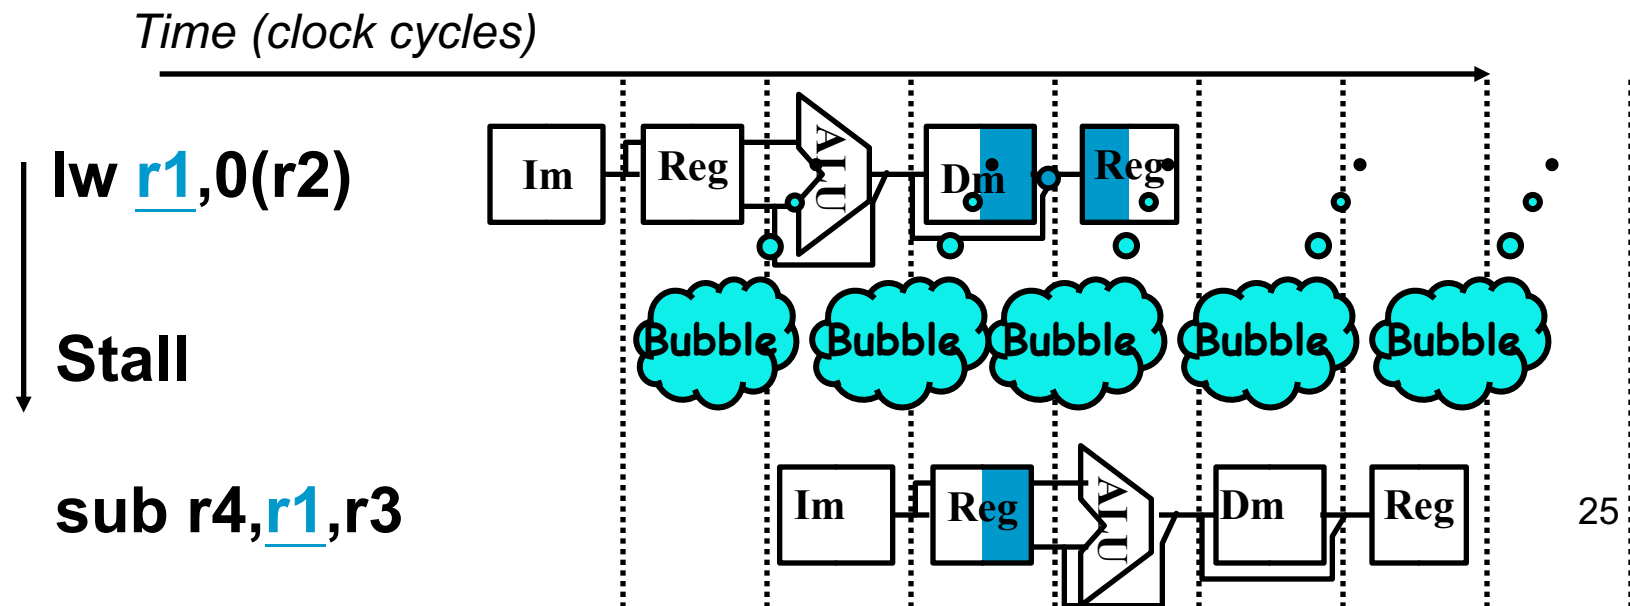


# Forwarding (or Bypassing): What about Loads

- **Load-use data hazard: cannot solve with forwarding**



- **Must delay/stall instruction dependent on loads**



# Reordering Code to Avoid Pipeline Stalls

---

A = B + E;



C = B + F;

```
lw      $t1, 0($t0)
lw      $t2, 4($t0)
add     $t3, $t1, $t2
sw      $t3, 12($t0)
lw      $t4, 8($t0)
add     $t5, $t1, $t4
sw      $t5, 16($t0)
```

# Reordering Code to Avoid Pipeline Stalls

A = B + E;  
C = B + F;

On a pipelined processor with forwarding, the reordered sequence will complete in two fewer cycles than left-hand side sequence.



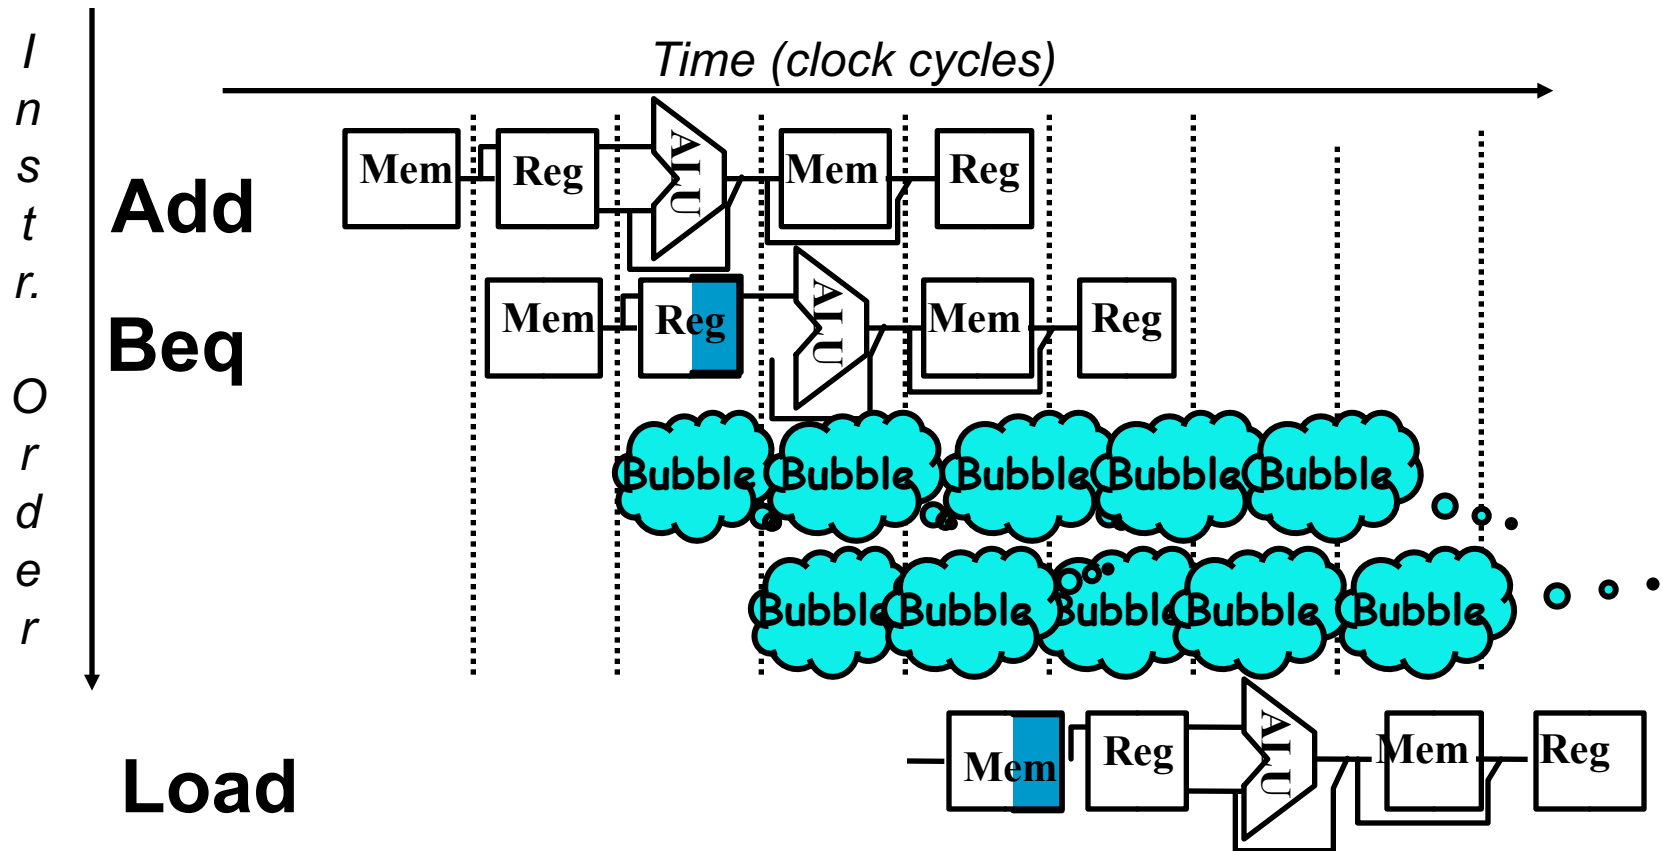
# Control Hazard

---

- It arises from the need to make a decision based on the results of one instruction while others are executing.
- Two solutions to control hazards
  - Stall
    - After fetching a branch instruction, pipeline is to stall immediately, waiting until the outcome of the branch is known.
    - Slow down the pipeline.
  - Predict
    - Assumes a given outcome for the branch and proceeds from the assumption rather than waiting to ascertain the actual outcome.
    - Should rollback if the assumption is wrong,

# Control Hazard Solutions (1) Stall

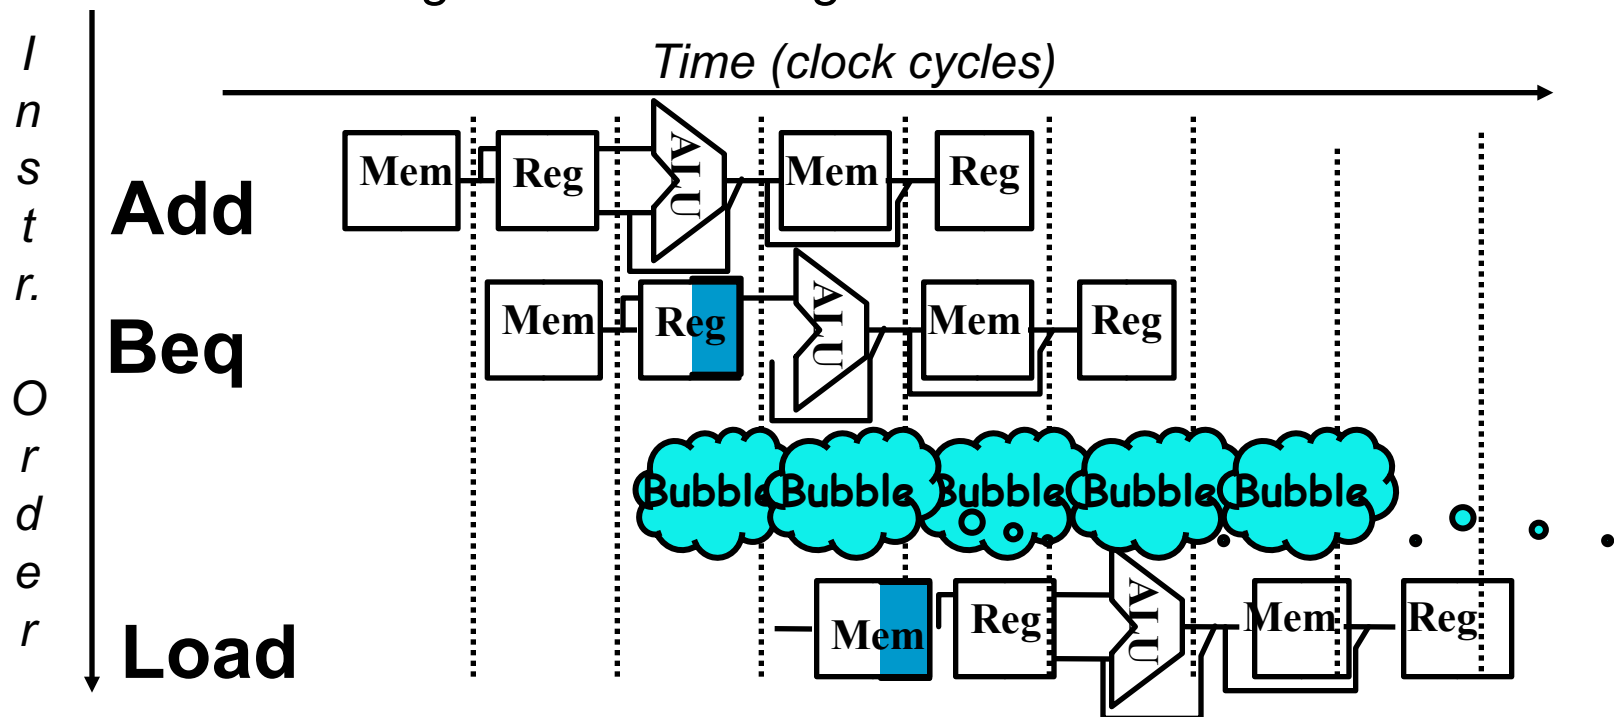
- When is a branch resolved?



- Impact: 3 clock cycles per branch instruction  
=> slow

## Control Hazard Solutions (2) Resolve branch earlier

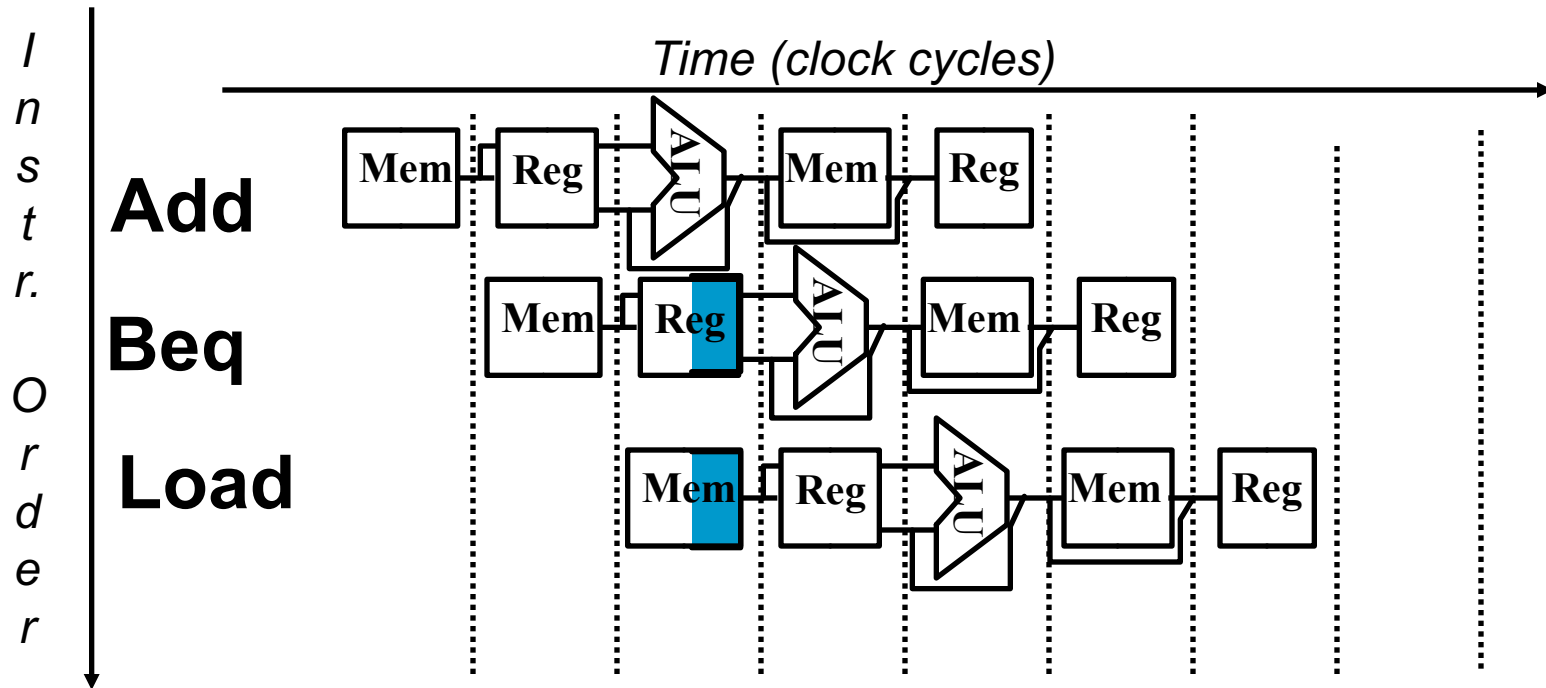
- What if a branch is resolved in the second stage (ID)?
  - Assume that we can put in enough extra hardware so that we can test registers, calculate the branch address, and update the PC during the second stage



- Impact: 2 clock cycles per branch instruction
- As the branch are 13% of the instructions executed in SPECint2000, the CPI is slowdown of 1.13 versus the ideal case.

# Control Hazard Solutions (3) Predicted Untaken

- Predict: guess one direction then back up if wrong
  - Predict not taken



- Impact: 1 clock cycles per branch instruction if right, 2 if wrong (right 50% of time)
- More **dynamic scheme**: history of 1 branch (90%)

## Predict branch not taken

---

Branch Inst (i)	IF	ID	EX	MEM	WB			
Inst i+1		IF	ID	EX	MEM	WB		
Inst i+2			IF	ID	EX	MEM	WB	
Inst i+3				IF	ID	EX	MEM	WB
Inst i+4					IF	ID	EX	MEM

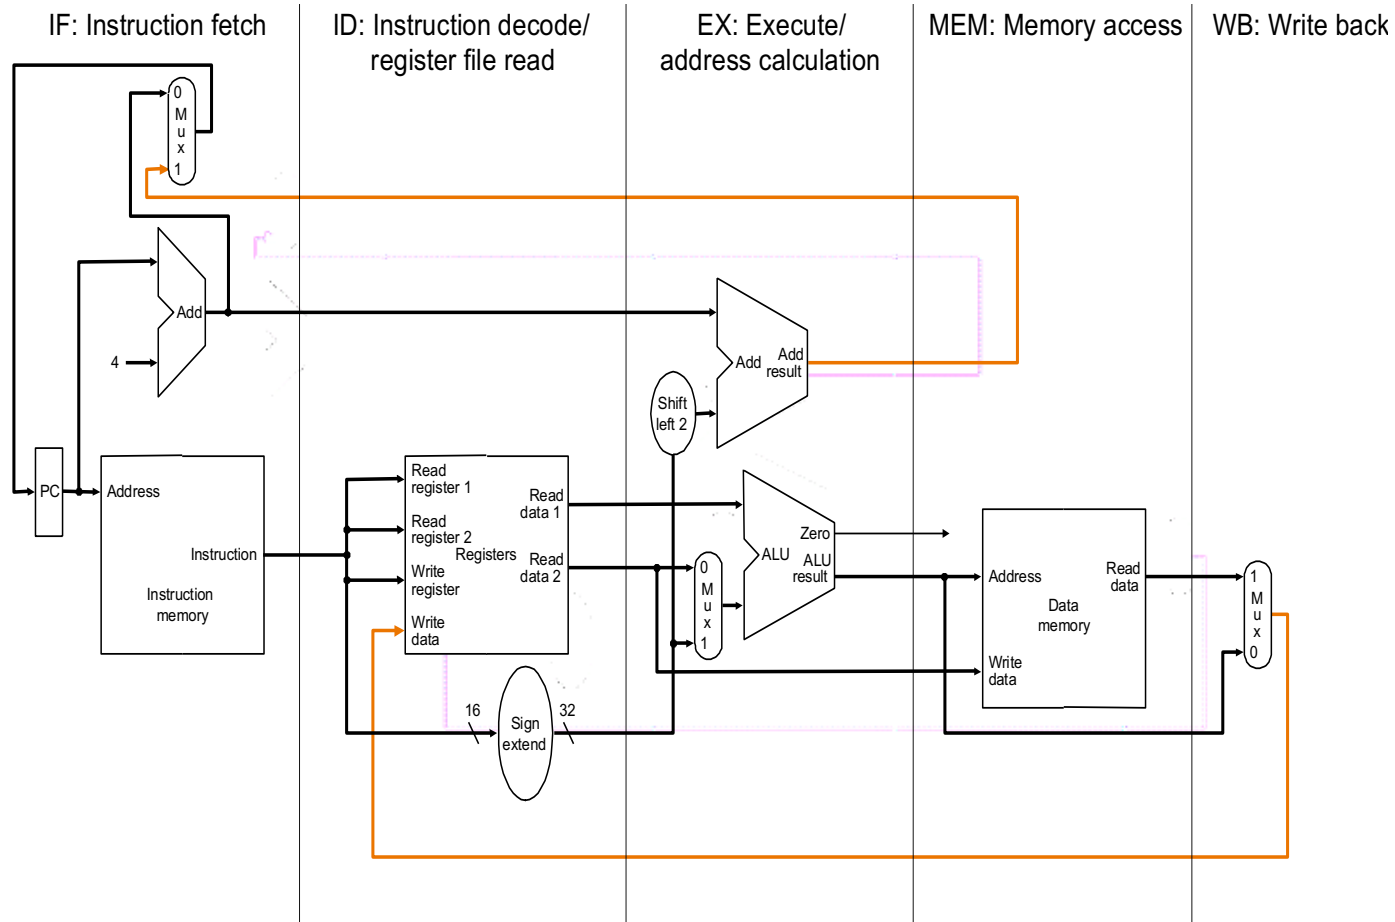
**Correct Prediction : Zero Cycle Branch Penalty!**

Branch Inst (i)	IF	ID	EX	MEM	WB			
Inst i+1		IF	nop	nop	nop	nop		
Branch target			IF	ID	EX	MEM	WB	

**Incorrect Prediction - waste one cycle**  
**How to flush pipeline?**

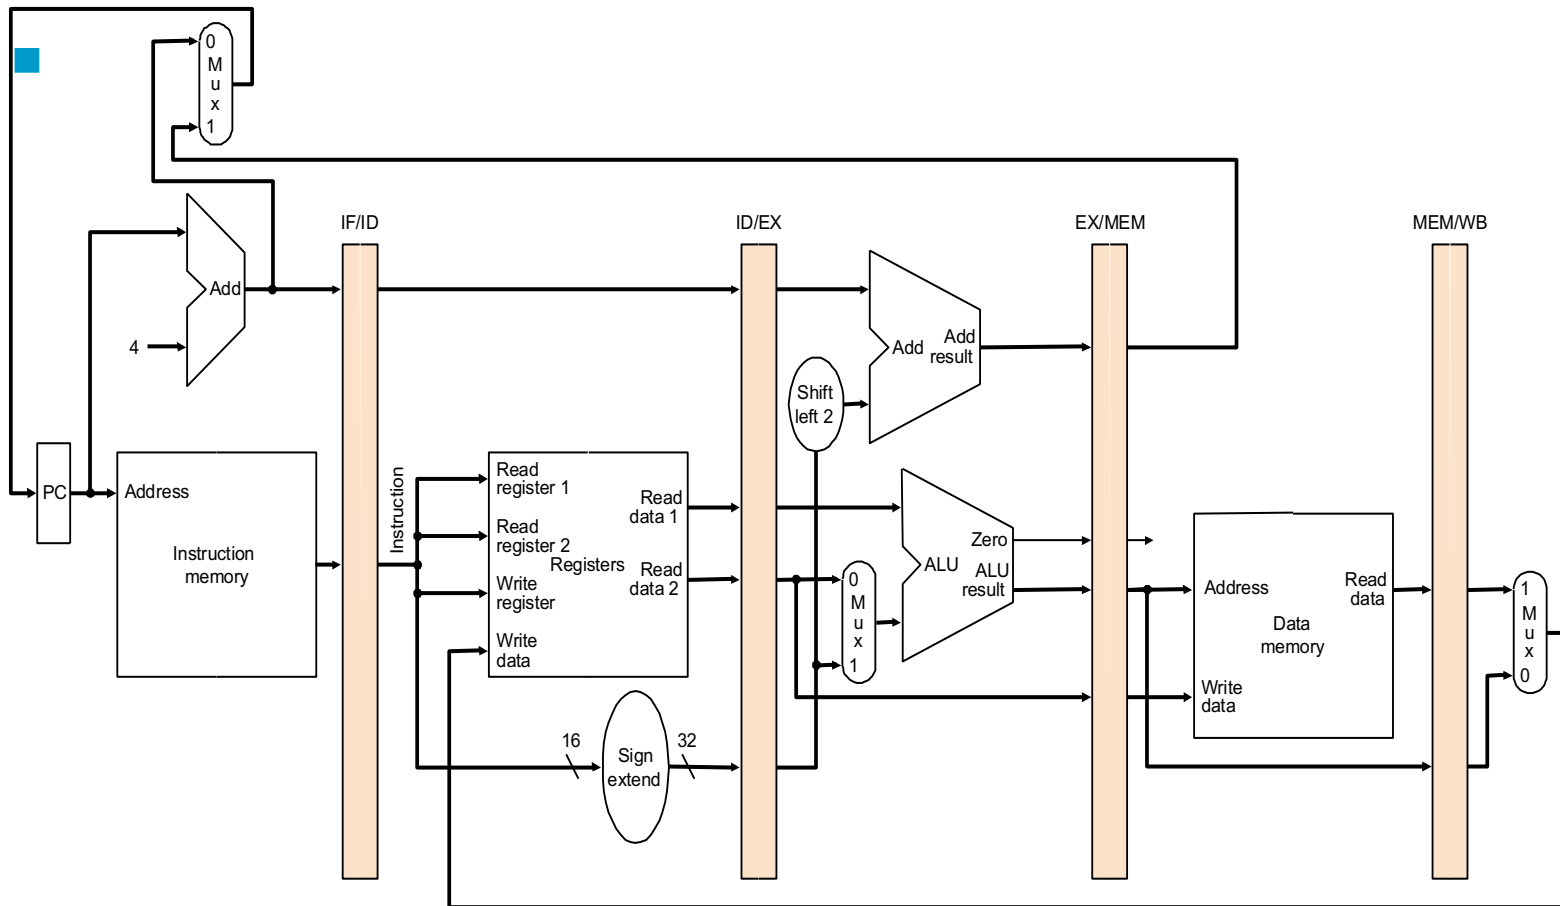


# A Pipelined Datapath: Basic Idea



■ What do we need to add to actually split the datapath into stages?

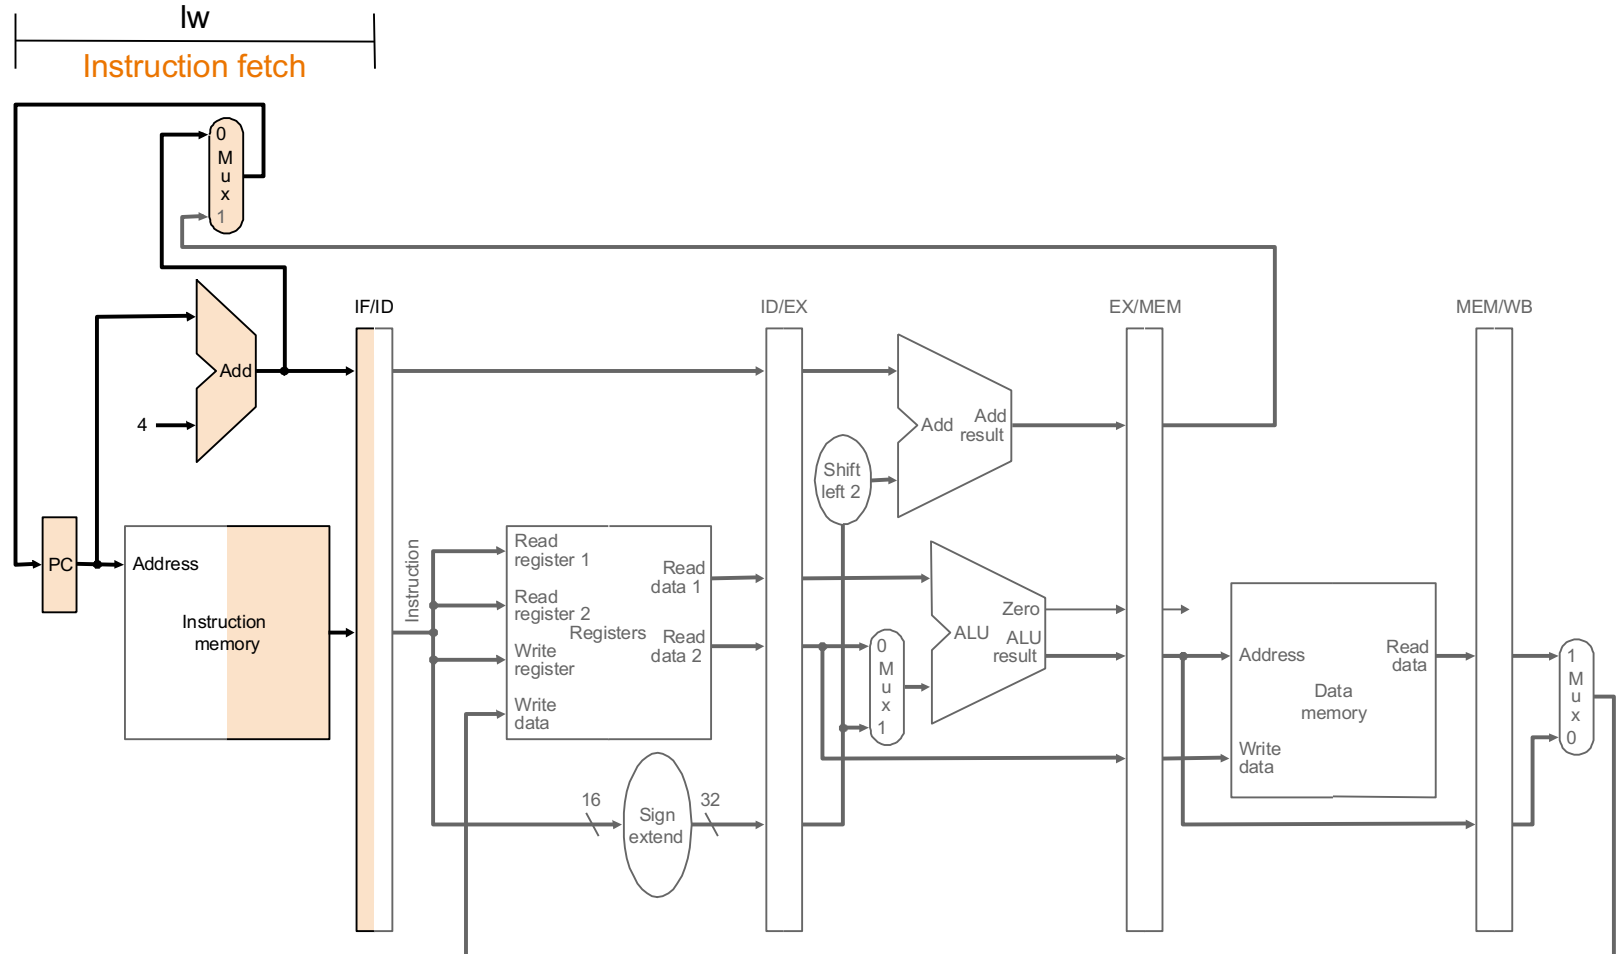
# Pipelined Datapath



- Pipeline registers:
  - hold data needed for later pipeline stage
  - The registers are named for the two stages separated by that register.
- Why do we need to have separate I/D data memory? Adder?

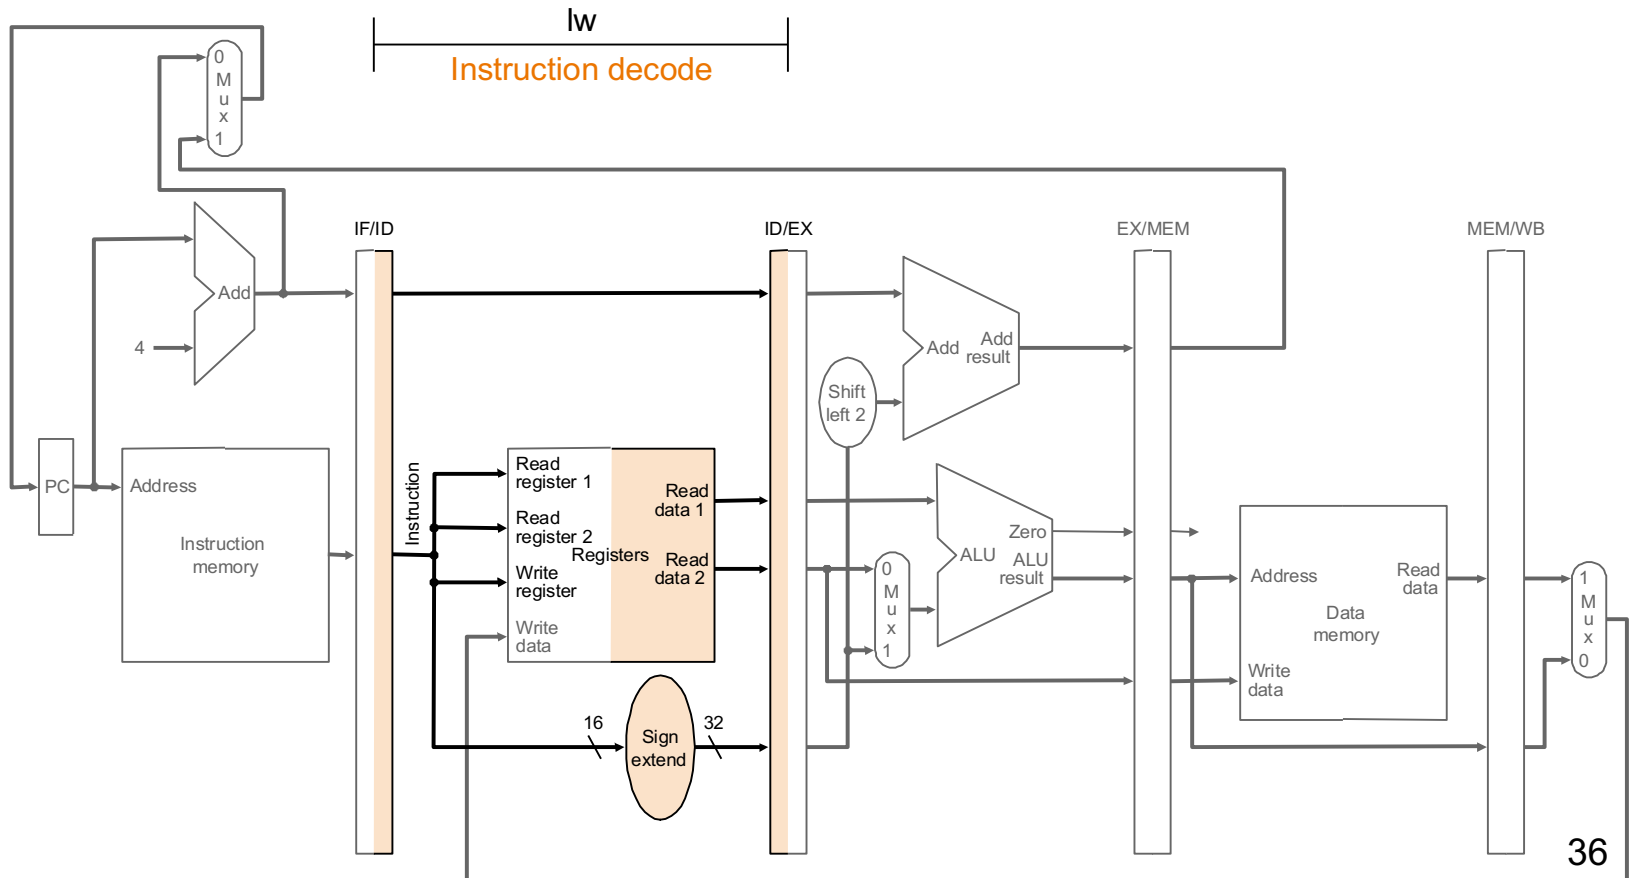
# IF Stage of load

- IF/ID.IR = mem[PC] ; IF/ID.PC = PC+4 ; PC = PC +4 ;



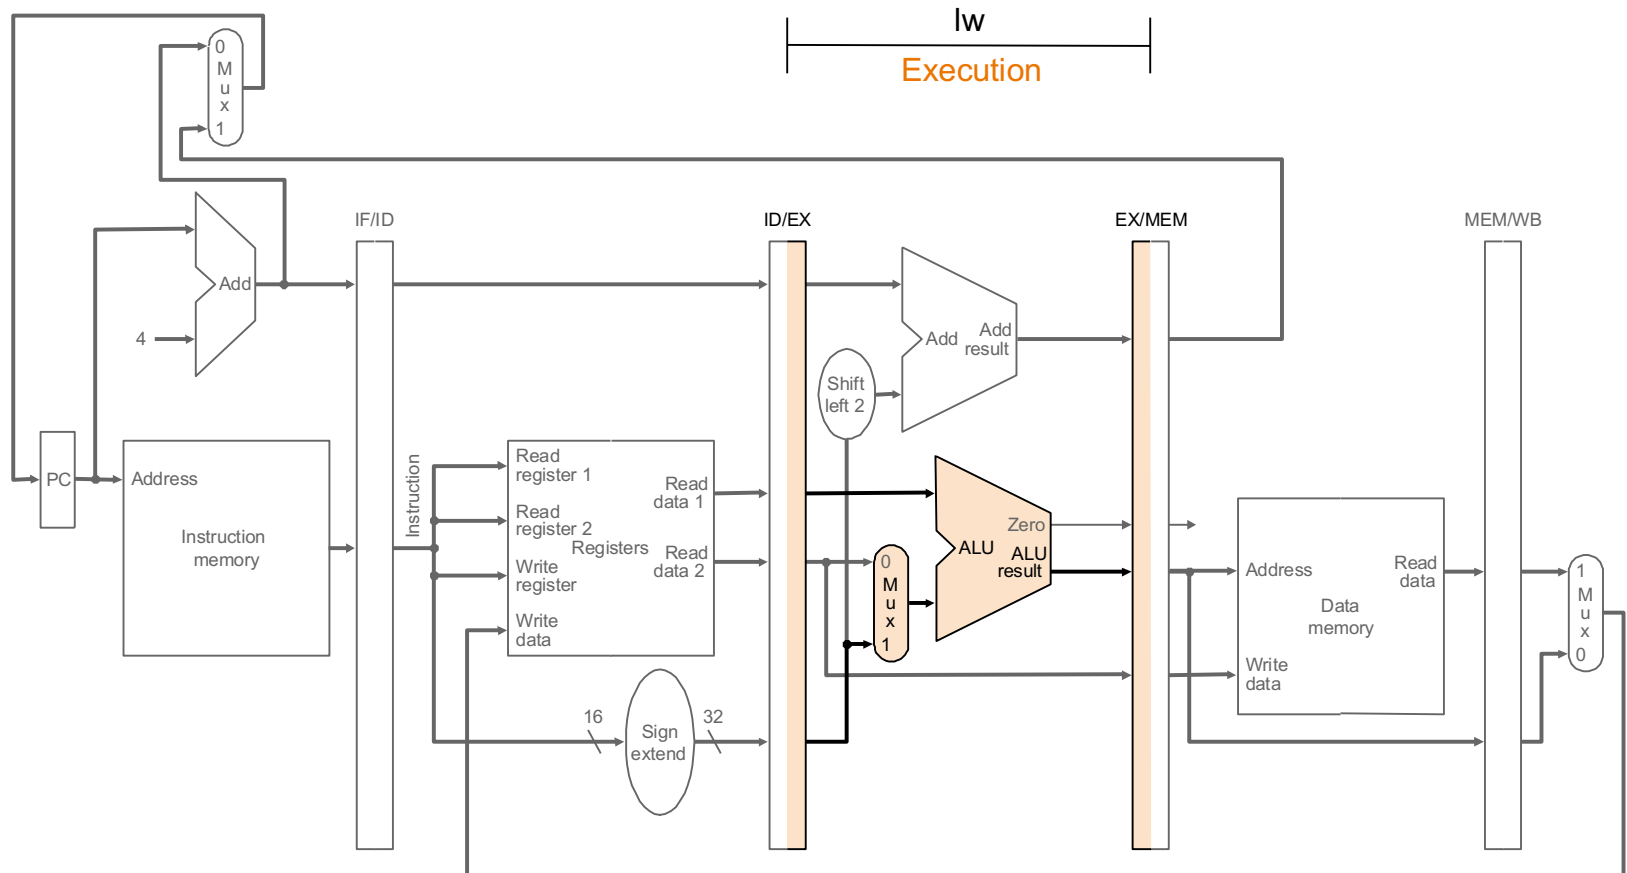
# ID Stage of load

- $ID/EX.A = Reg[IF/ID.IR[25-21]]$
- $ID/EX.B = Reg[IF/ID.IR[20-16]]$
- $ID/EX.Immediate = (sign-ext(IF/ID.IR[15-0]))$
- $ID/EX.PC = IF/ID.PC$



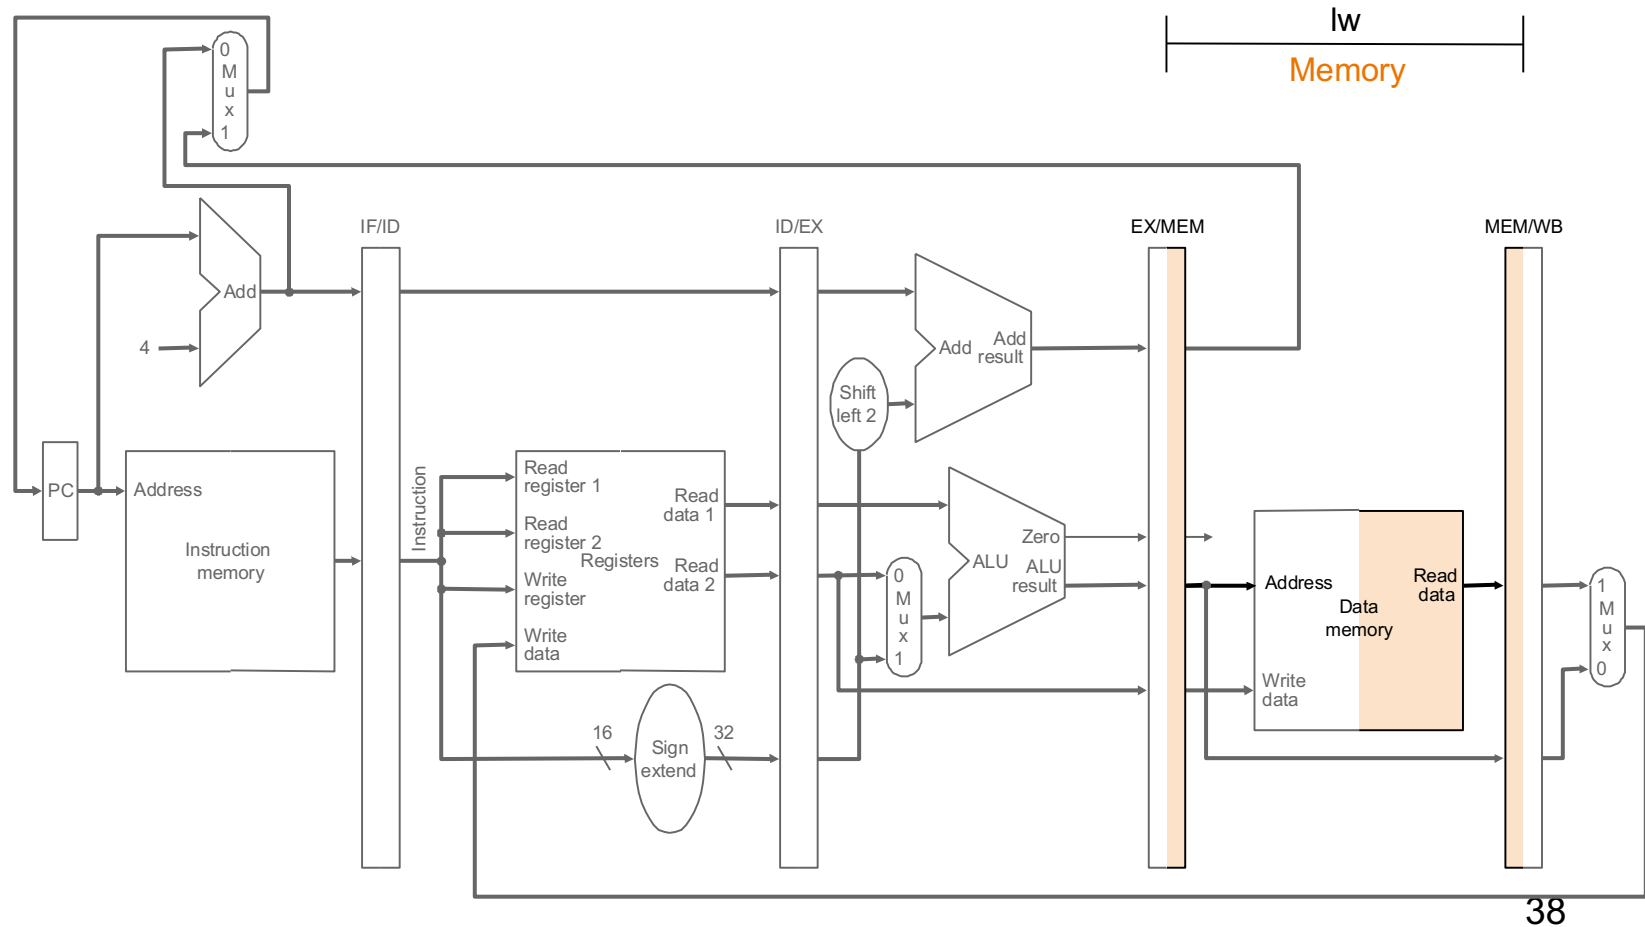
## EX Stage of load

- $EX/MEM.ALUout = ID/EX.A + ID/EX.immediate$



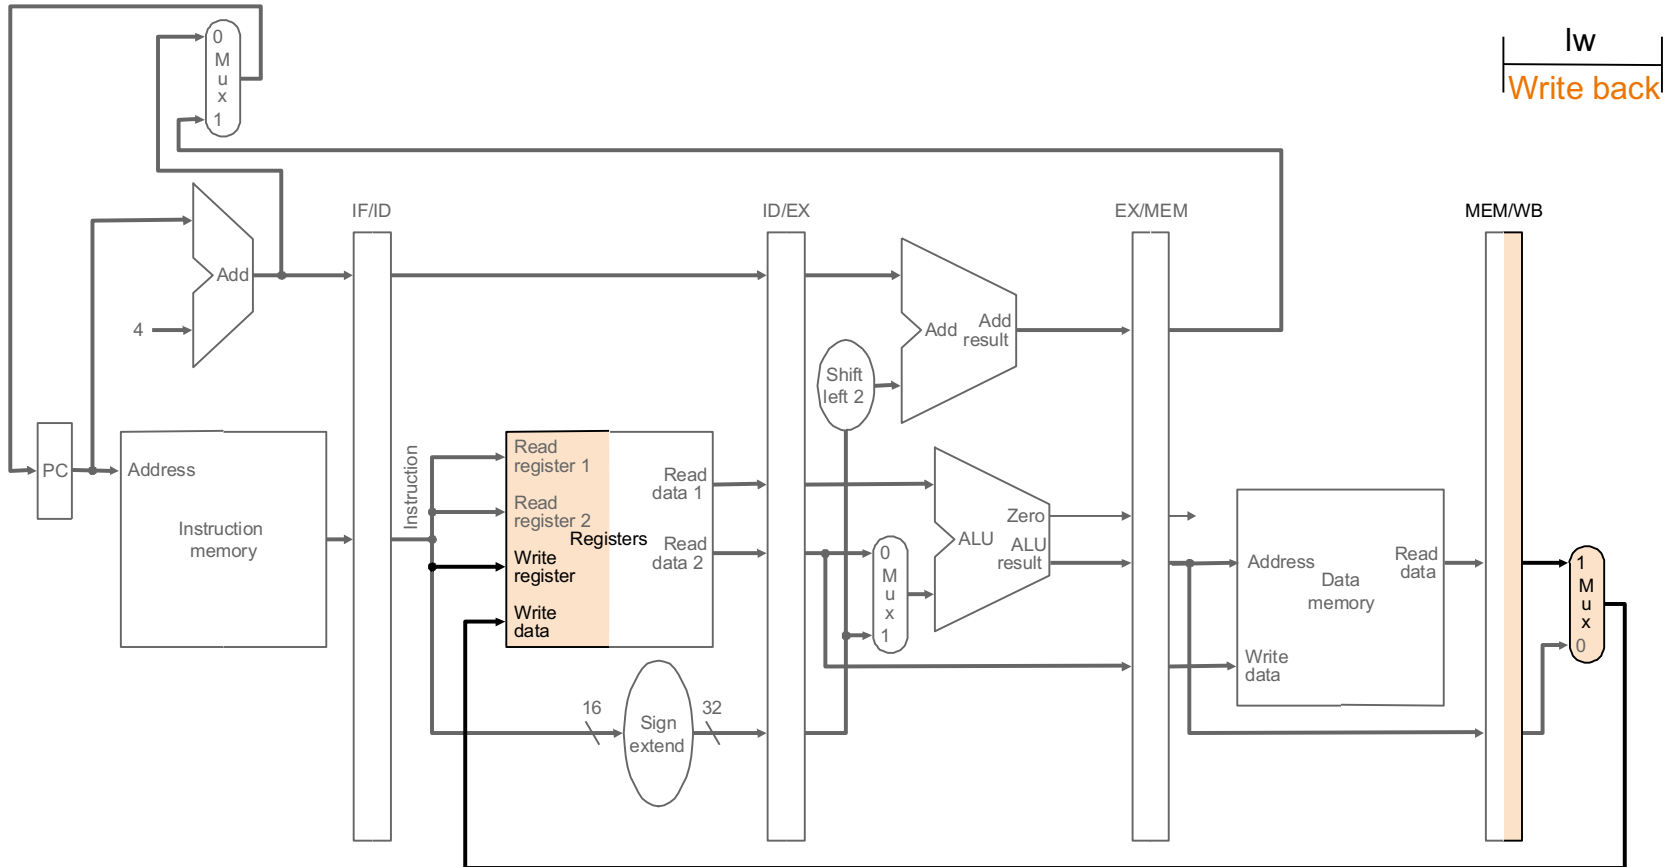
## MEM Stage of load

- MEM/WB.MDR = mem[EX/MEM.ALUout]



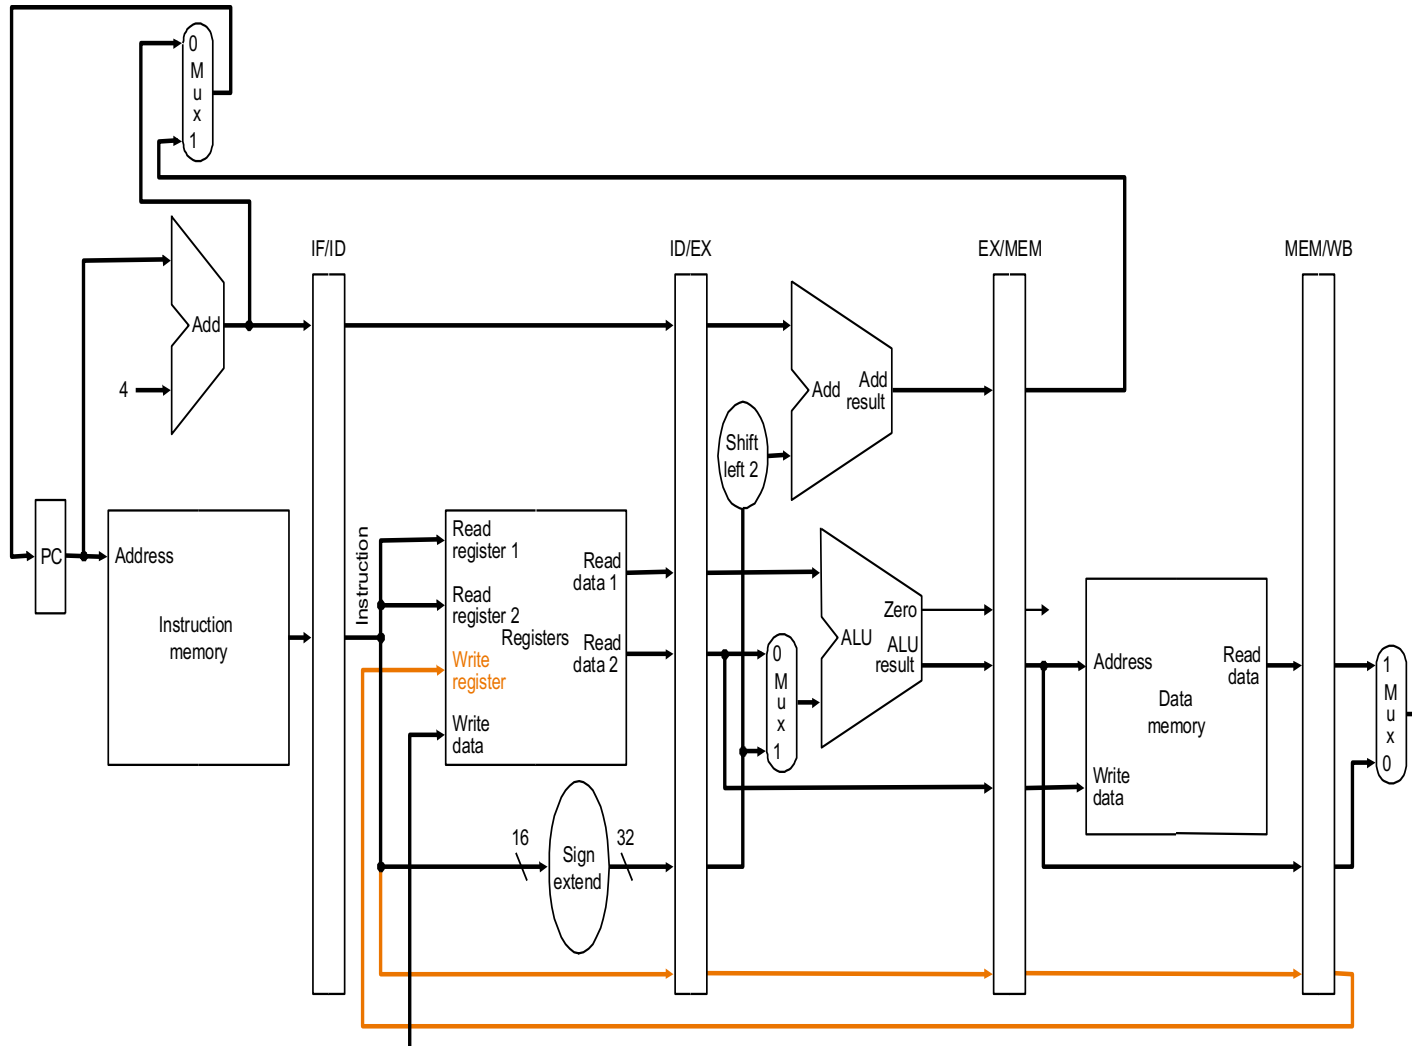
# WB Stage of load

- $\text{Reg}[\text{Rt}] = \text{MEM/WB.MDR}$ 
  - Where does Rt come from?



- Load must pass the register number from ID/EX through EX/MEM to the MEM/WB pipeline register for use in the WB stage

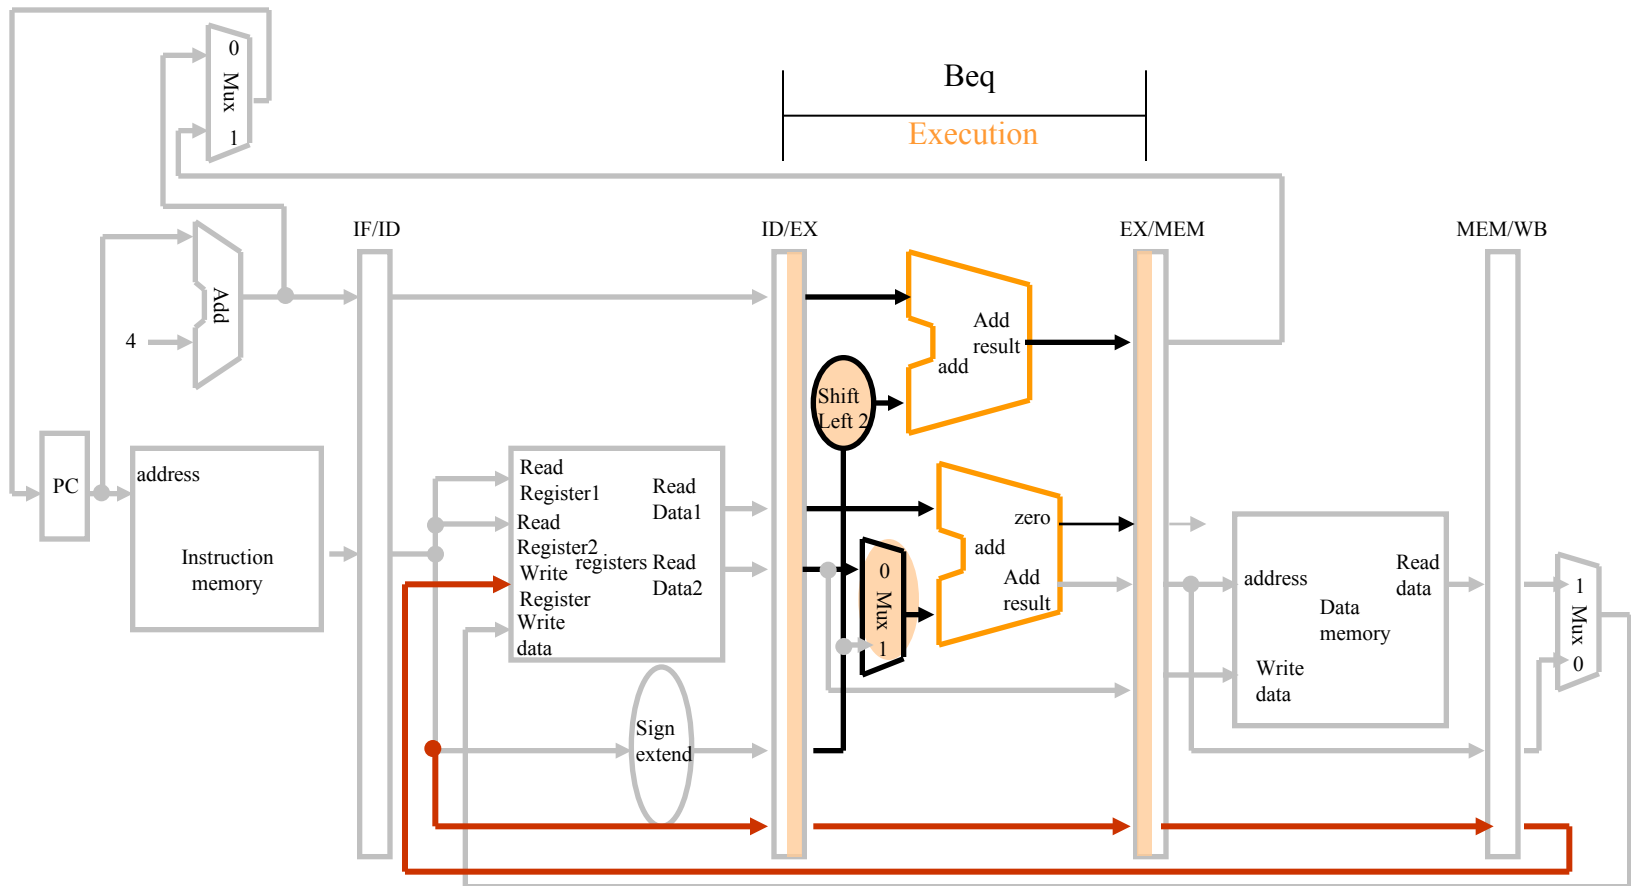
# Corrected Datapath





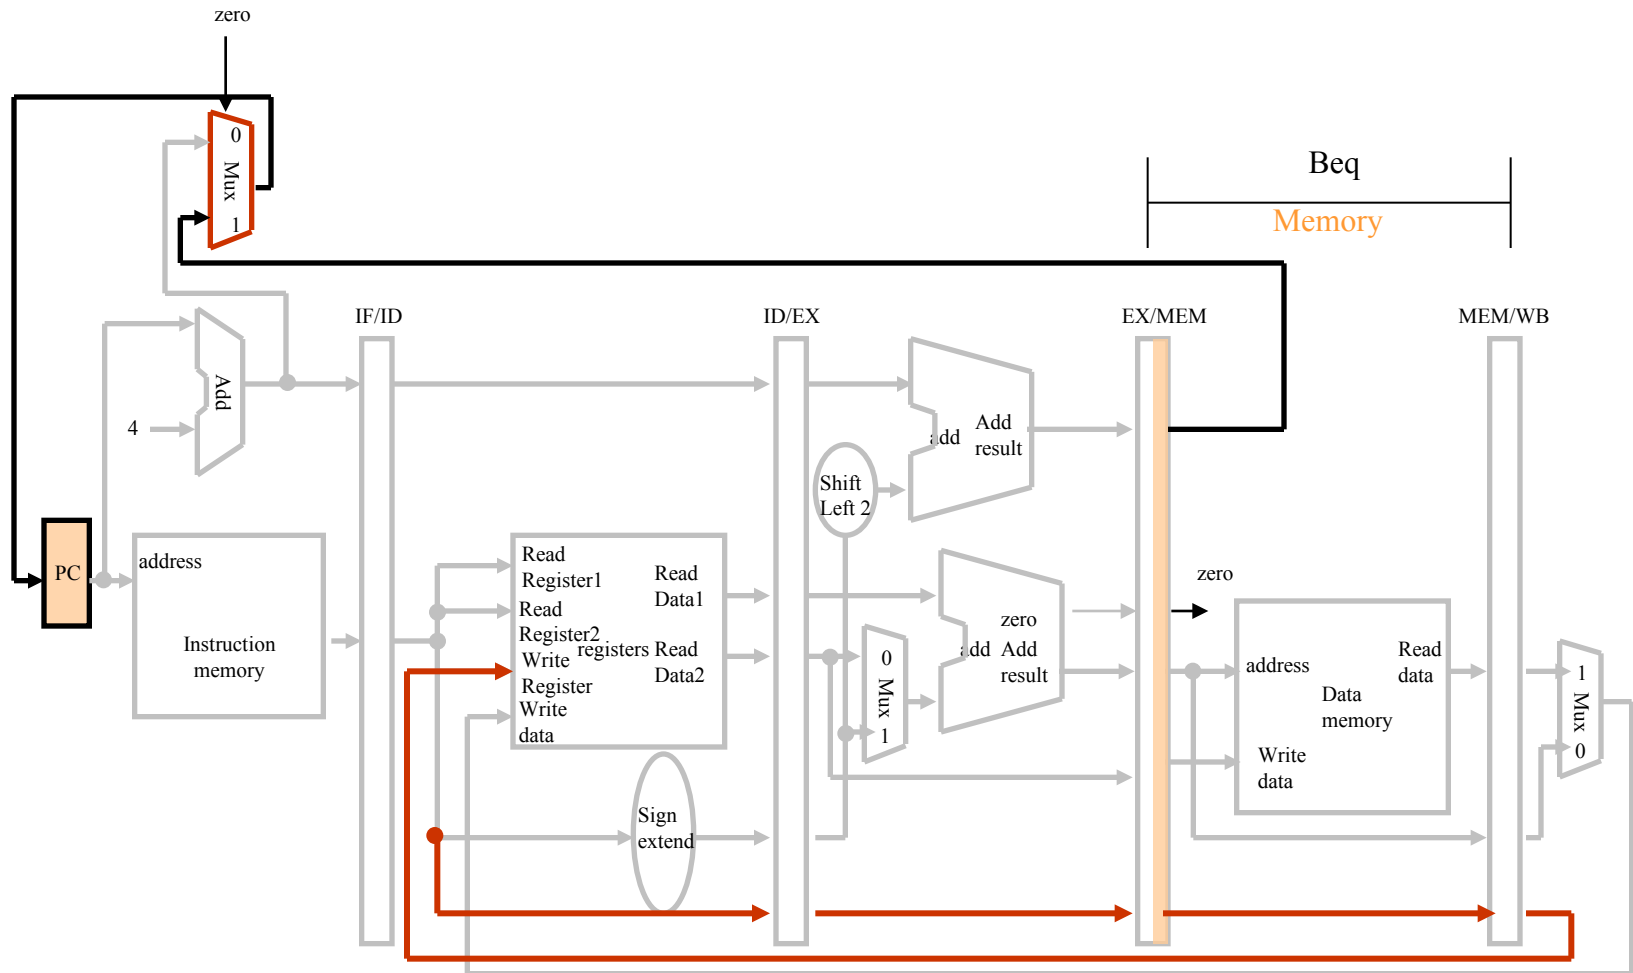
# EX Stage of beq

- $EX/MEM.PC = ID/EX.PC + (ID/EX.immediate \ll 2);$
- $EX/MEM.Zero = (ID/EX.A - ID/EX.B) ? 0 : 1$

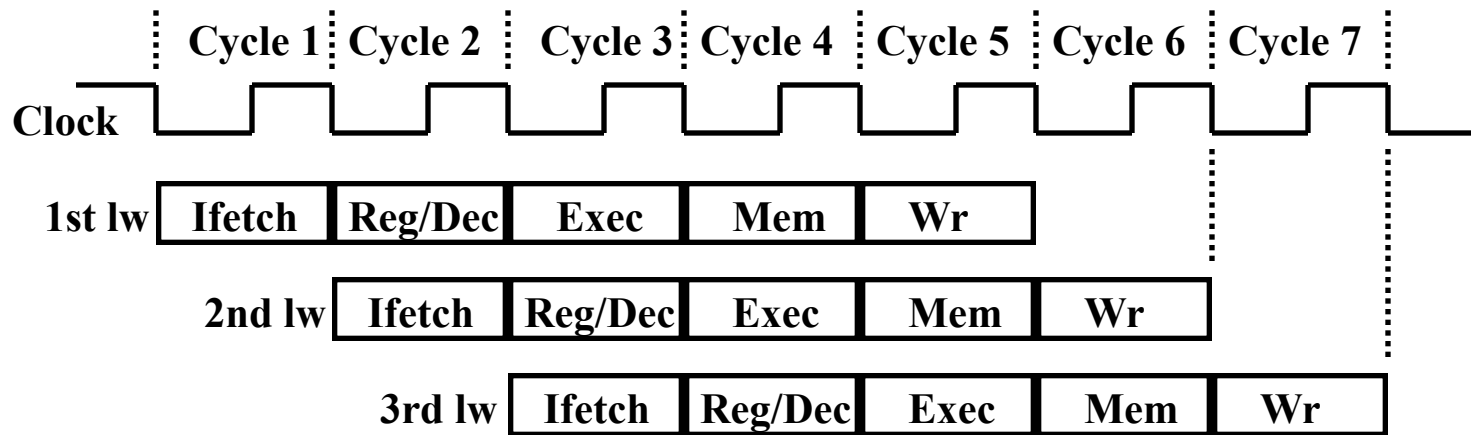


# MEM Stage of beq

- If (EX/MEM.zero) PC = EX/MEM.PC




# Pipelining the Load Instruction



- The five independent functional units in the pipeline datapath are:
  - Instruction Memory for the **Ifetch** stage
  - Register File's Read ports (bus A and busB) for the **Reg/Dec** stage
  - ALU for the **Exec** stage
  - Data Memory for the **Mem** stage
  - Register File's **Write** port (bus W) for the **Wr** stage

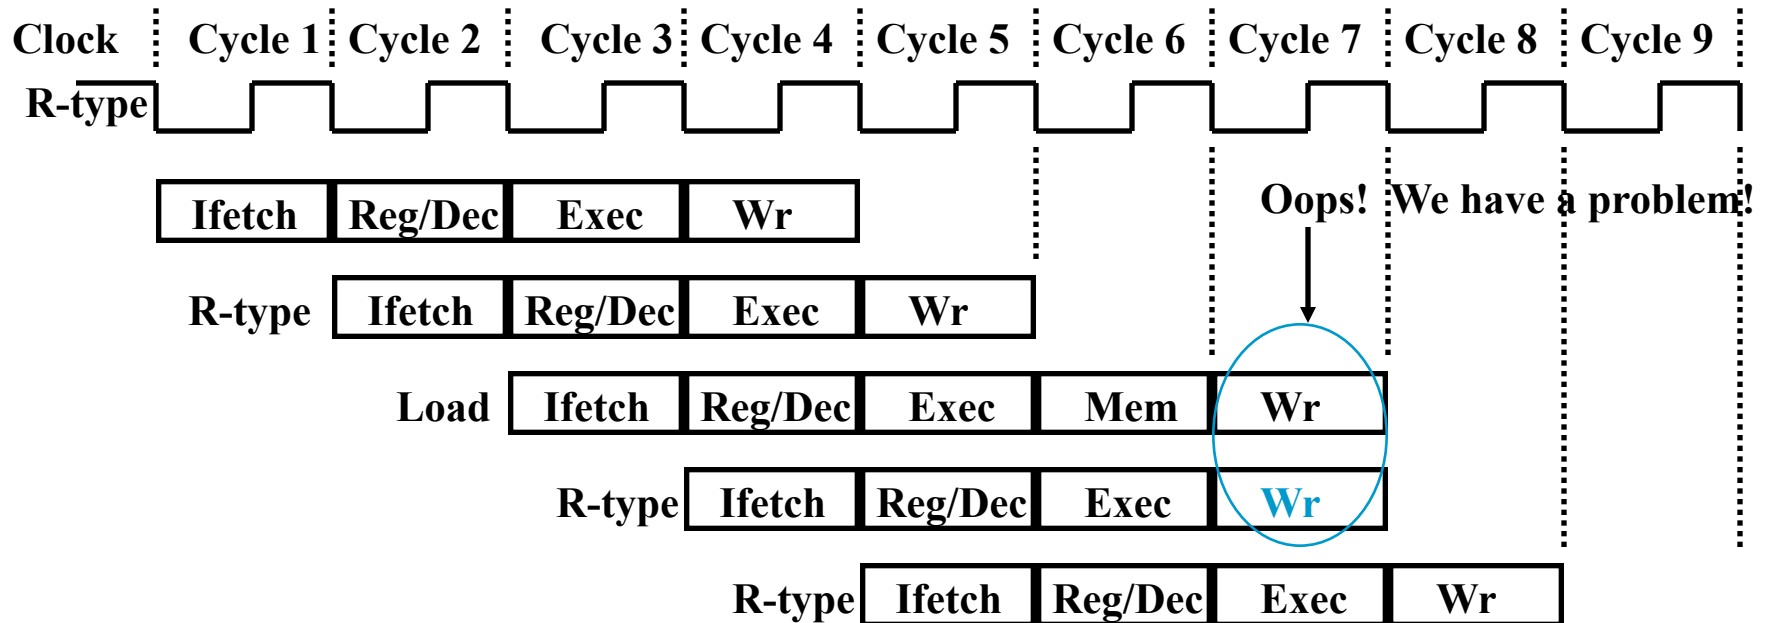
# The Four Stages of R-type

---



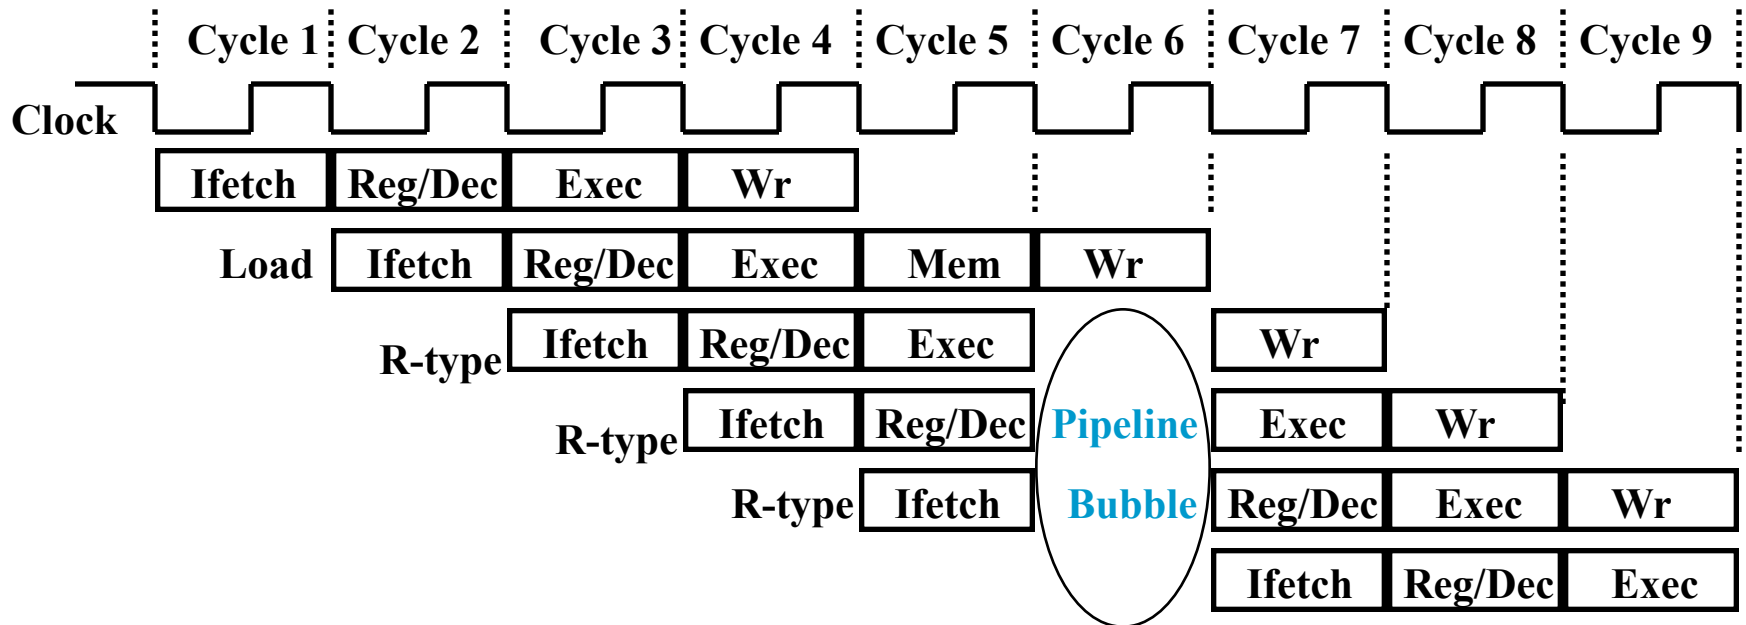
- **Ifetch: Instruction Fetch**
  - Fetch the instruction from the Instruction Memory
- **Reg/Dec: Registers Fetch and Instruction Decode**
- **Exec:**
  - ALU operates on the two register operands
- **Wr: Write the ALU output back to the register file**

# Pipelining the R-type and Load Instruction



- We have pipelined conflict or structural hazard:
  - Two instructions try to write to the register file at the same time!
  - Only one write port

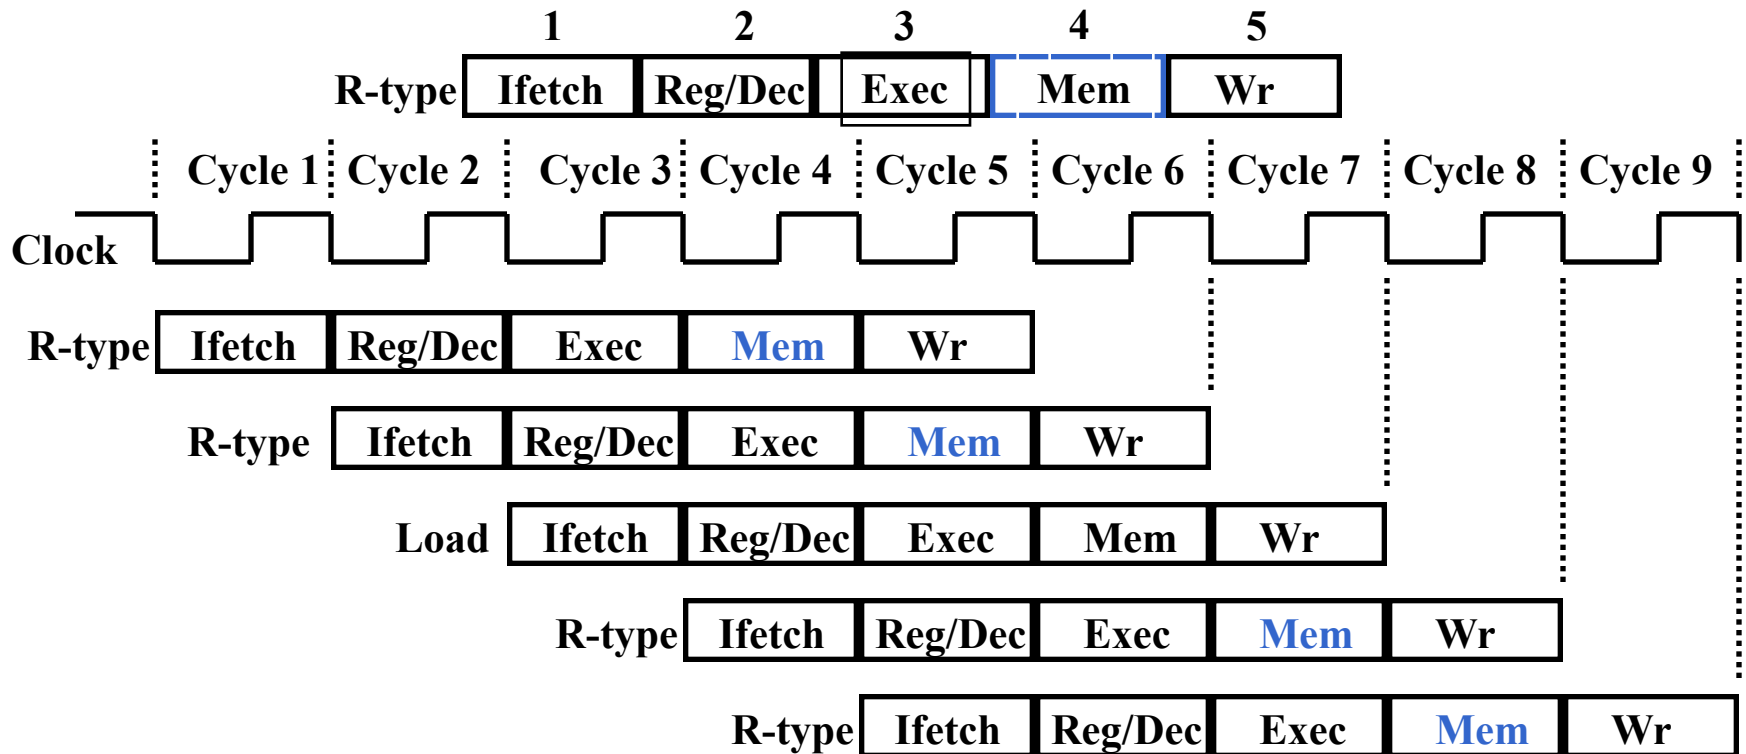
# Solution 1: Insert “Bubble” into the Pipeline



- Insert a “bubble” into the pipeline to prevent 2 writes at the same cycle
  - The control logic can be complex.
  - Lose instruction fetch and issue opportunity.
    - No instruction is started in Cycle 6!


## Solution 2: Delay R-type's Write by One Cycle

- Delay R-type's register write by one cycle:
  - Now R-type instructions also use Reg File's write port at Stage 5
  - Mem stage is a **NOOP** stage: nothing is being done.



# The Four Stages of Store

---

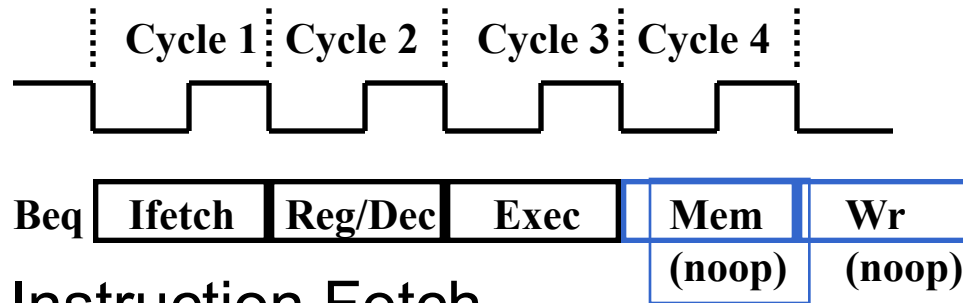


- Ifetch: Instruction Fetch
  - Fetch the instruction from the Instruction Memory
- Reg/Dec: Registers Fetch and Instruction Decode
- Exec: Calculate the memory address
- Mem: Write the data into the Data Memory



# The Three Stages of Beq

---



## ■ Ifetch: Instruction Fetch

- Fetch the instruction from the Instruction Memory

## ■ Reg/Dec:

- Registers Fetch and Instruction Decode

## ■ Exec:

- compares the two register operand,
- select correct branch target address
- latch into PC

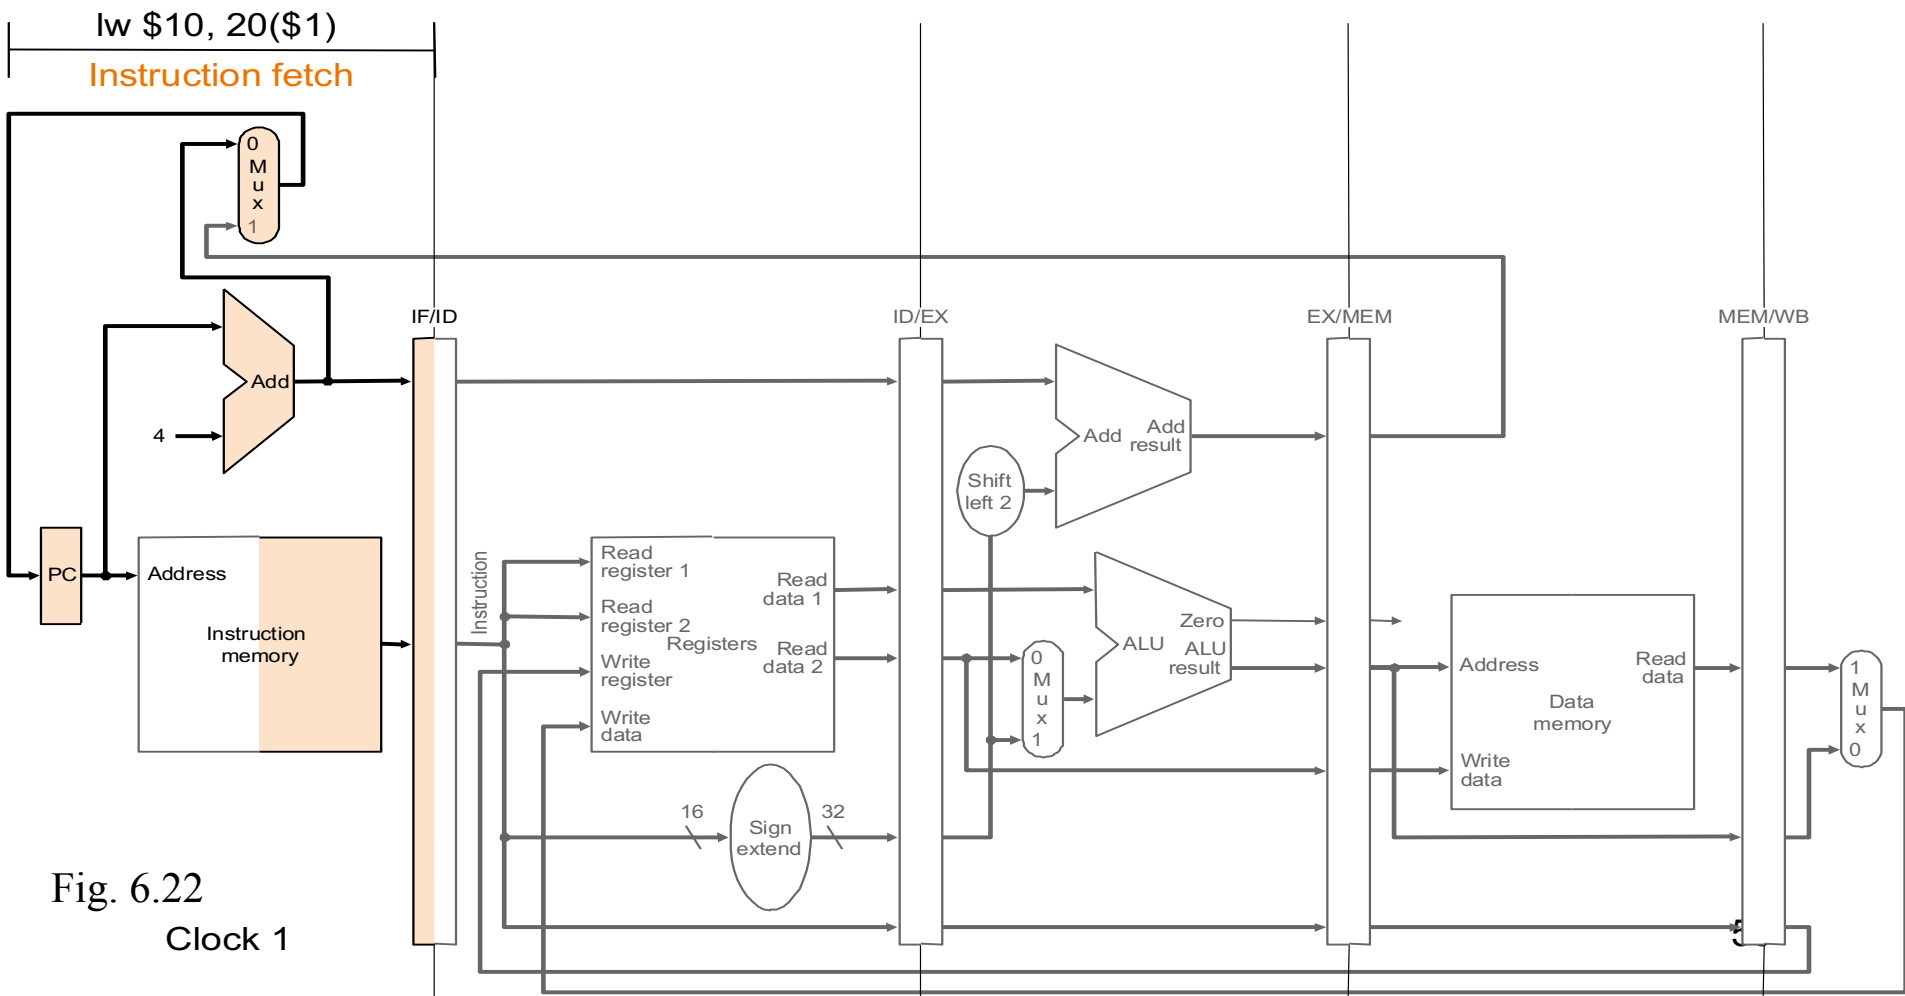


Fig. 6.22  
Clock 1





Fig. 6.22






Fig. 6.23

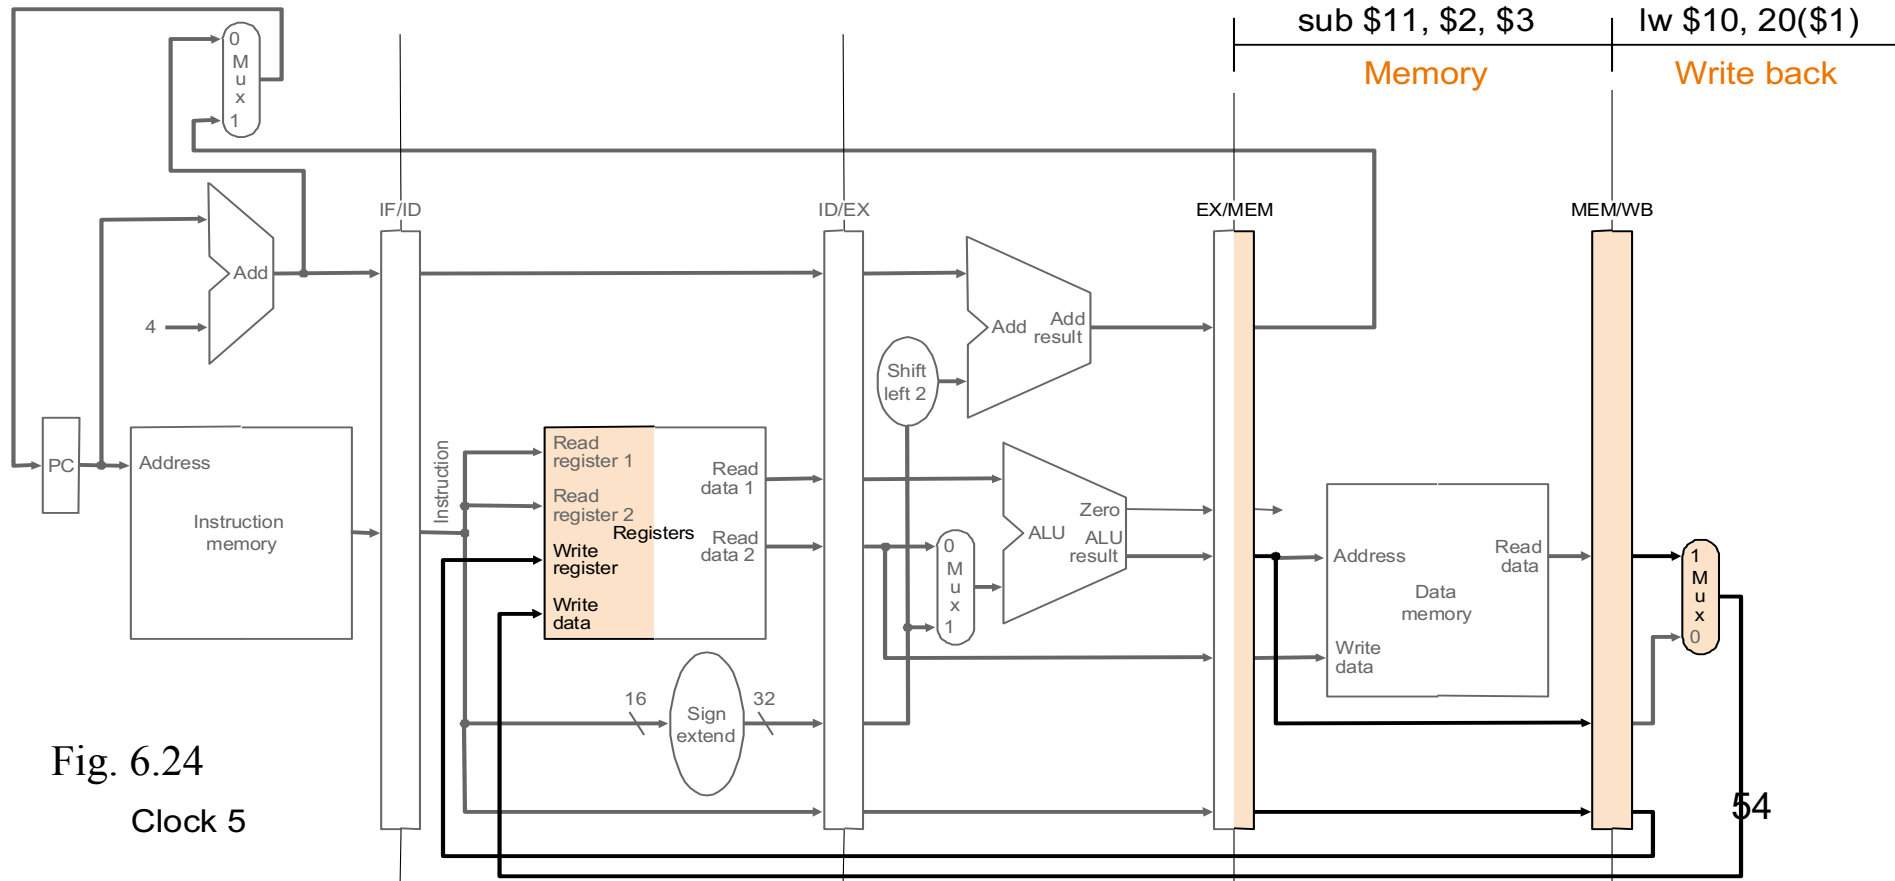


Fig. 6.24  
Clock 5

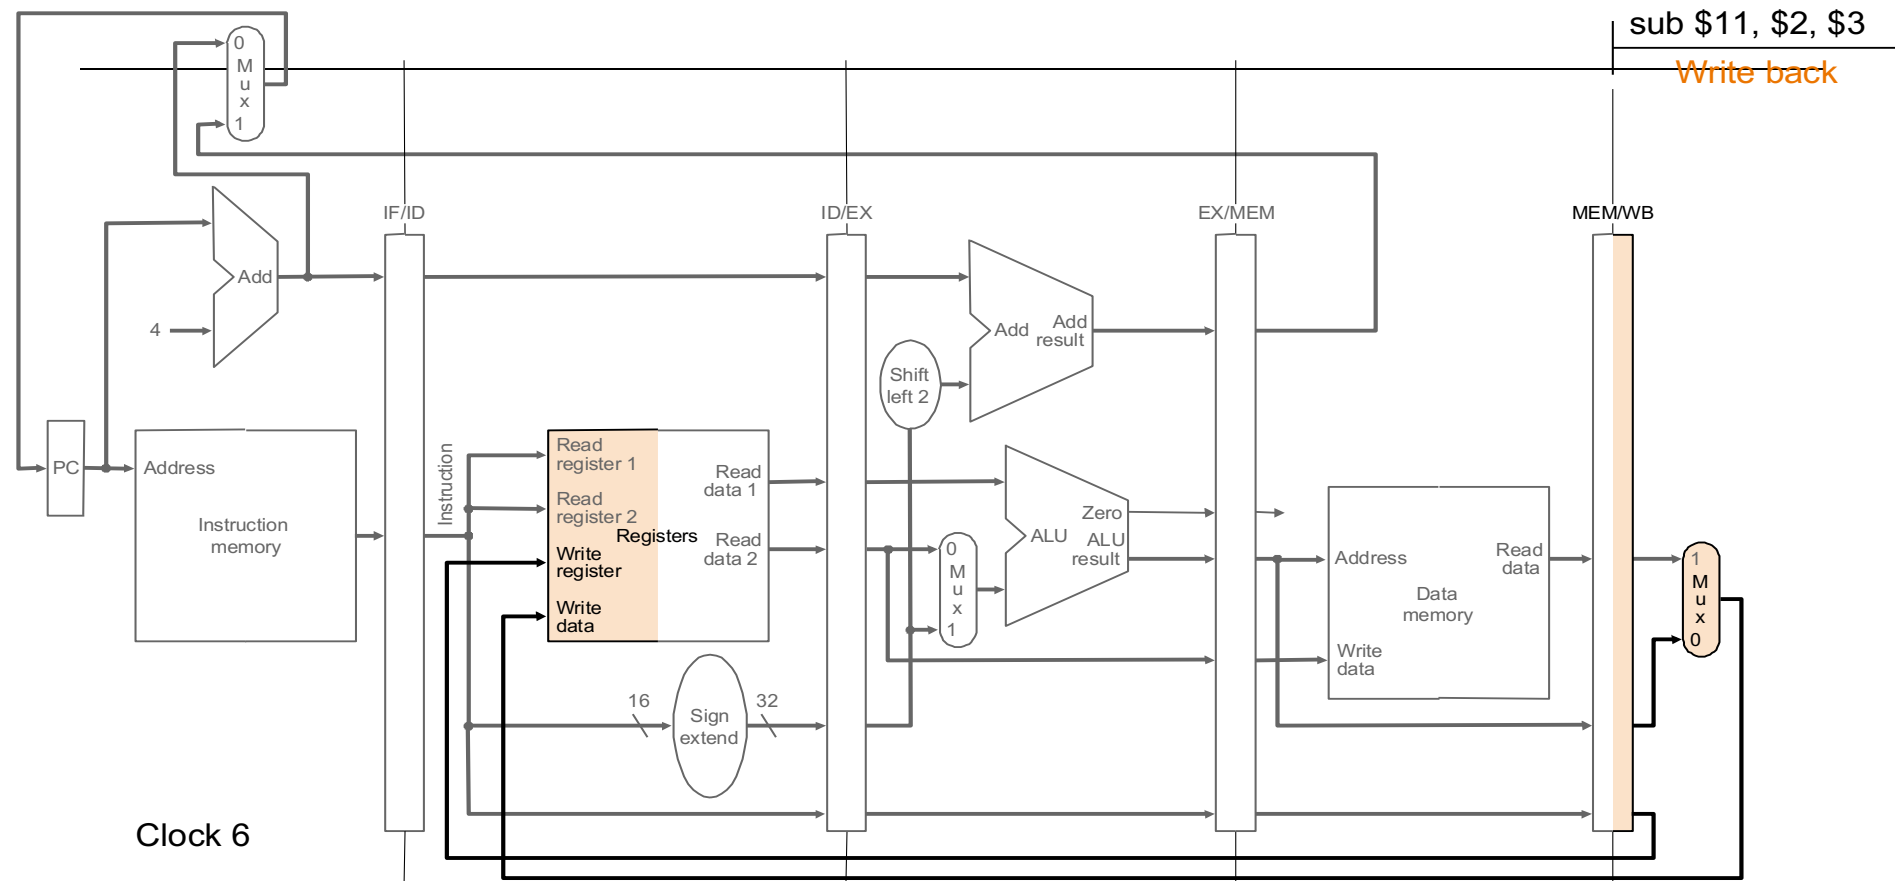


Fig. 6.24

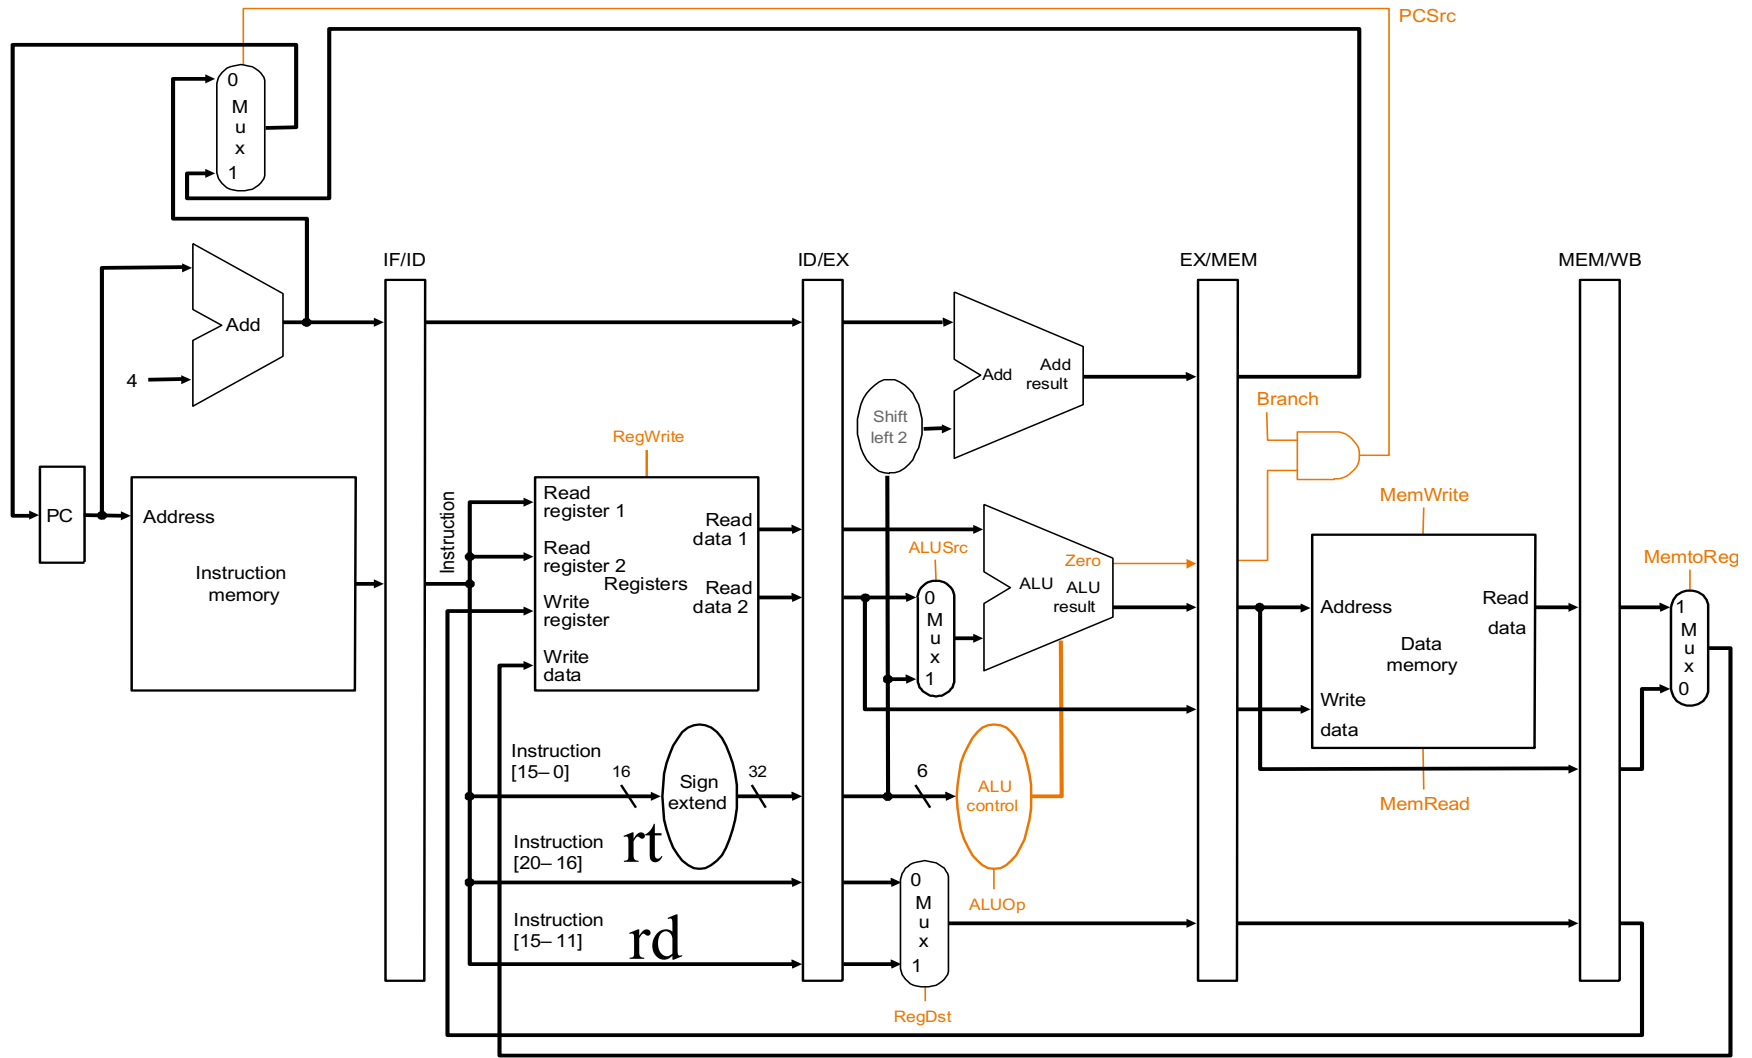
# Pipeline control

---

- We have 5 stages. What needs to be controlled in each stage?
  - Instruction Fetch and PC Increment
  - Instruction Decode / Register Fetch
  - Execution
  - Memory Stage
  - Write Back



# Pipelined Datapath with Control Signals Identified



1. Program counter is updated in the mem stage for beq
2. Do we need PCwrite control signal?


# Group Signals According to Stages

- Can use control signals of single-cycle CPU (Fig. 6.23, 6.24, 6.25  $\Leftrightarrow$  5.12, 5.16, 5.18)

Execution/Address Calculation stage control lines				Memory access stage control lines			Write-back stage control lines	
Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-type	1	0	0	0	0	0	1	0
lw	0	0	1	0	1	0	1	1
sw	X	0	1	0	0	1	0	X
Beq	X	1	0	1	0	0	0	X

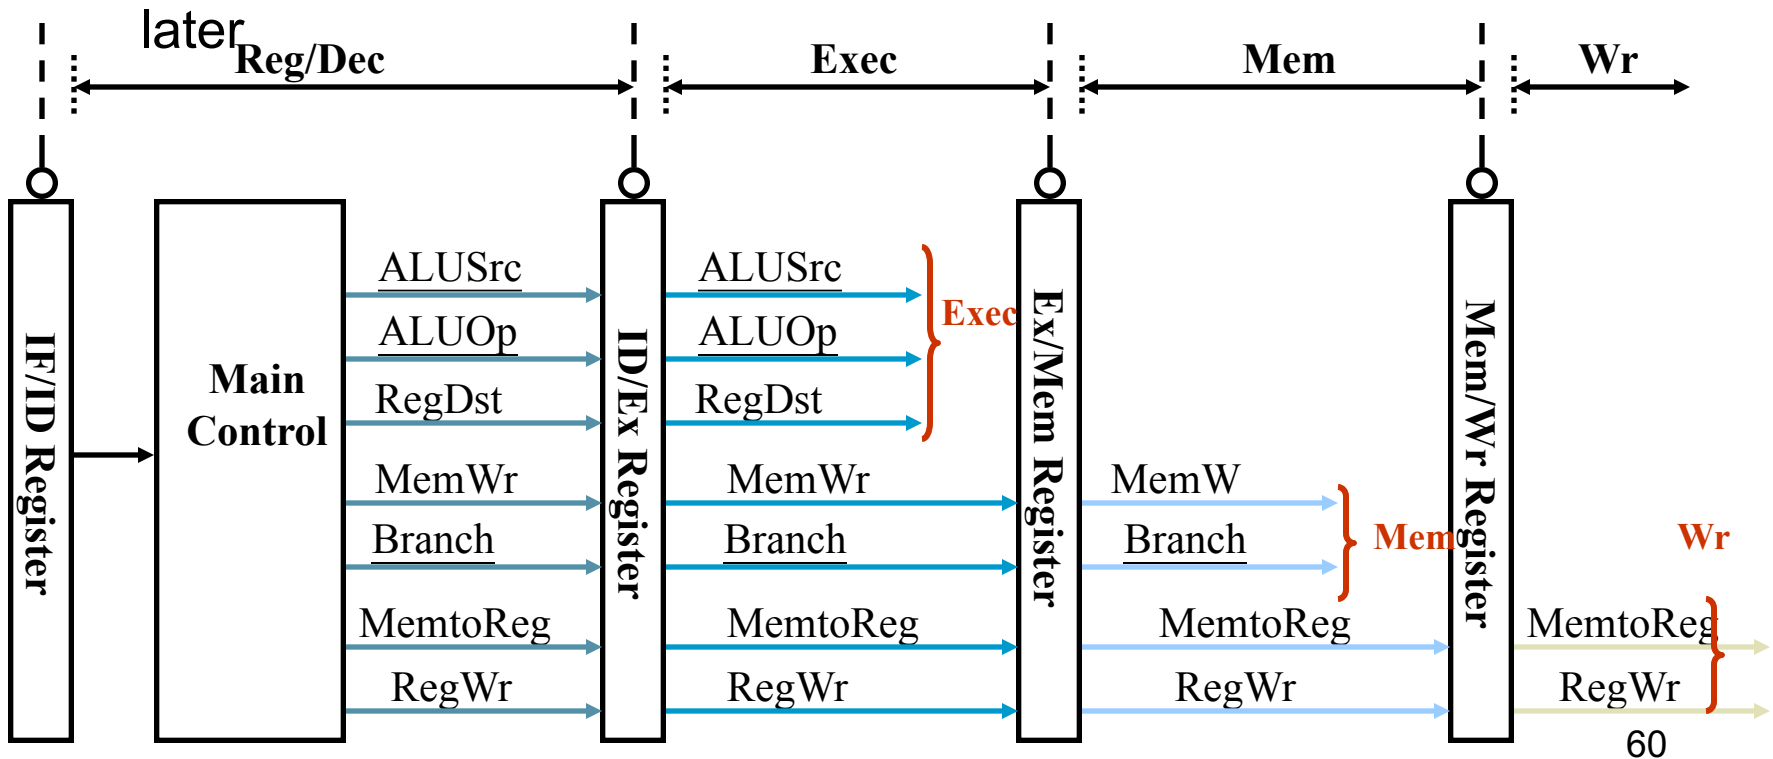
# Control Lines for the Final Three Stages

- Pass control signals along just like the data
  - Main control generates control signals during ID

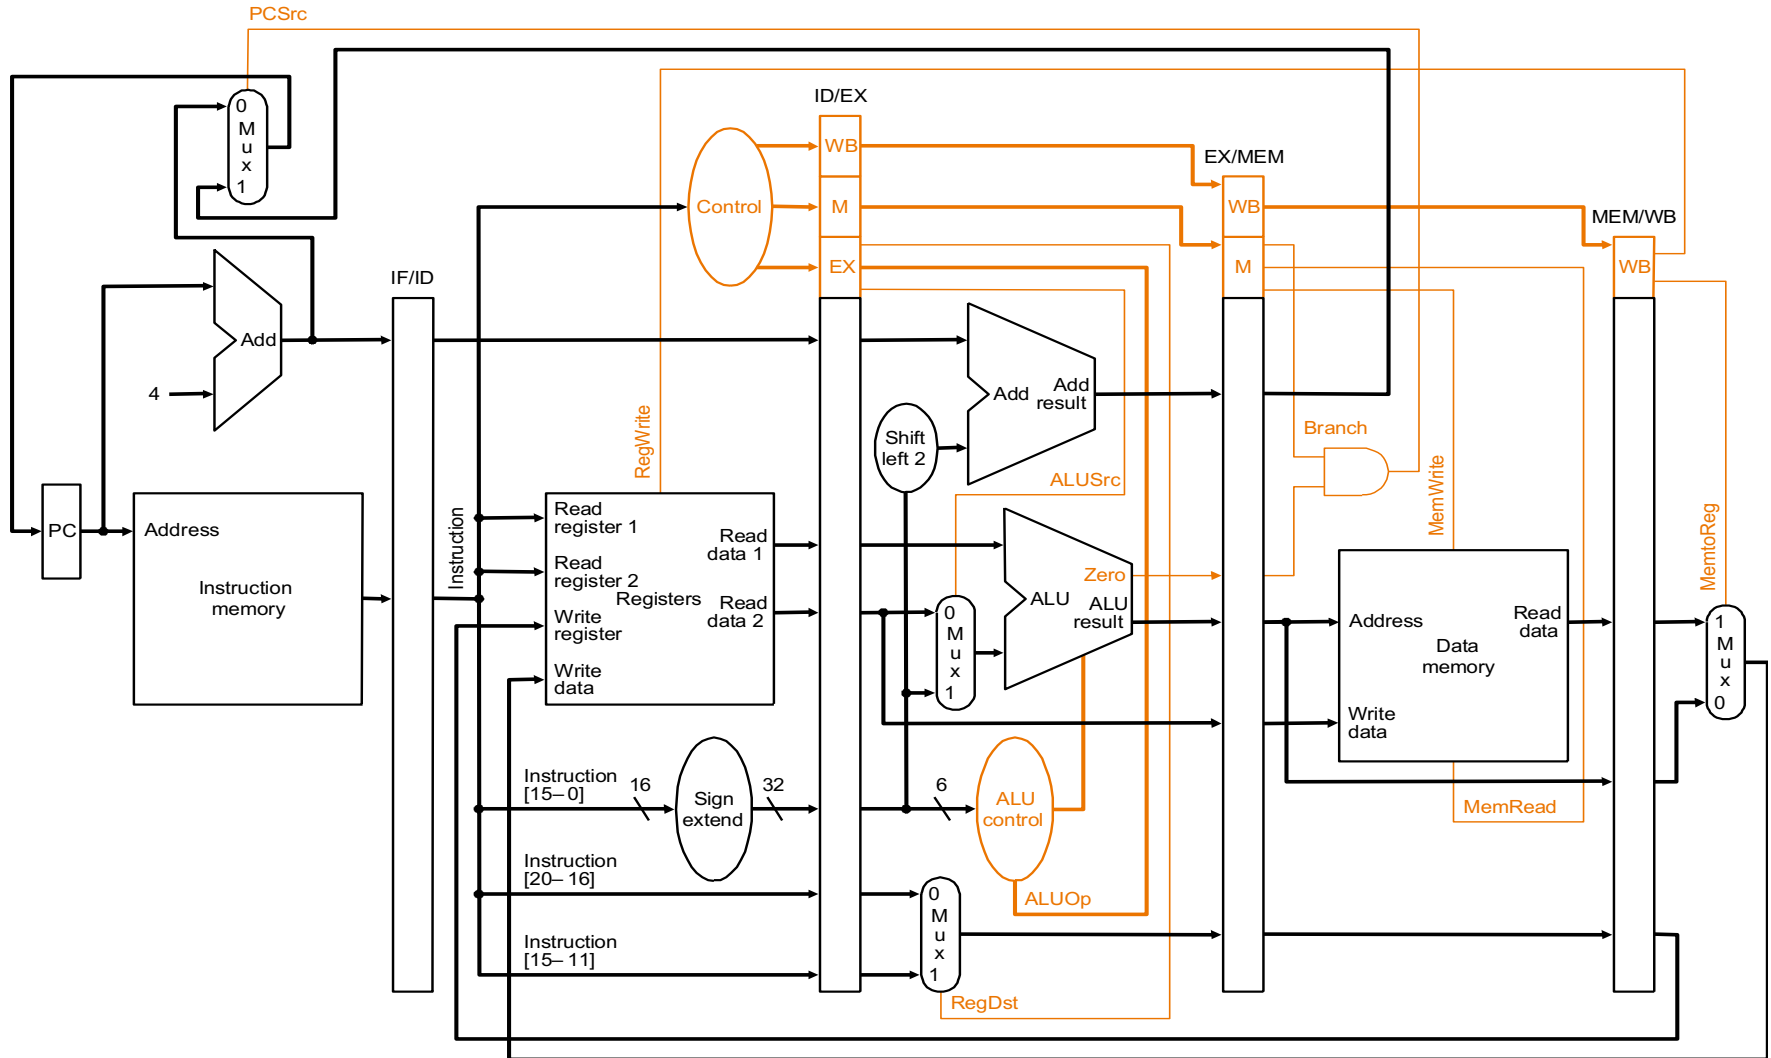


# Pipeline Control

- The Main Control generates the control signals during Reg/Dec
  - Control signals for Exec (ALUSrc, ...) are used 1 cycle later
  - Control signals for Mem (MemWr Branch) are used 2 cycles later
  - Control signals for Wr (MemtoReg MemWr) are used 3 cycles later



# Datapath with Control



# Instructions Pipelining Example

---

**lw**      **\$10, 20 (\$1)**

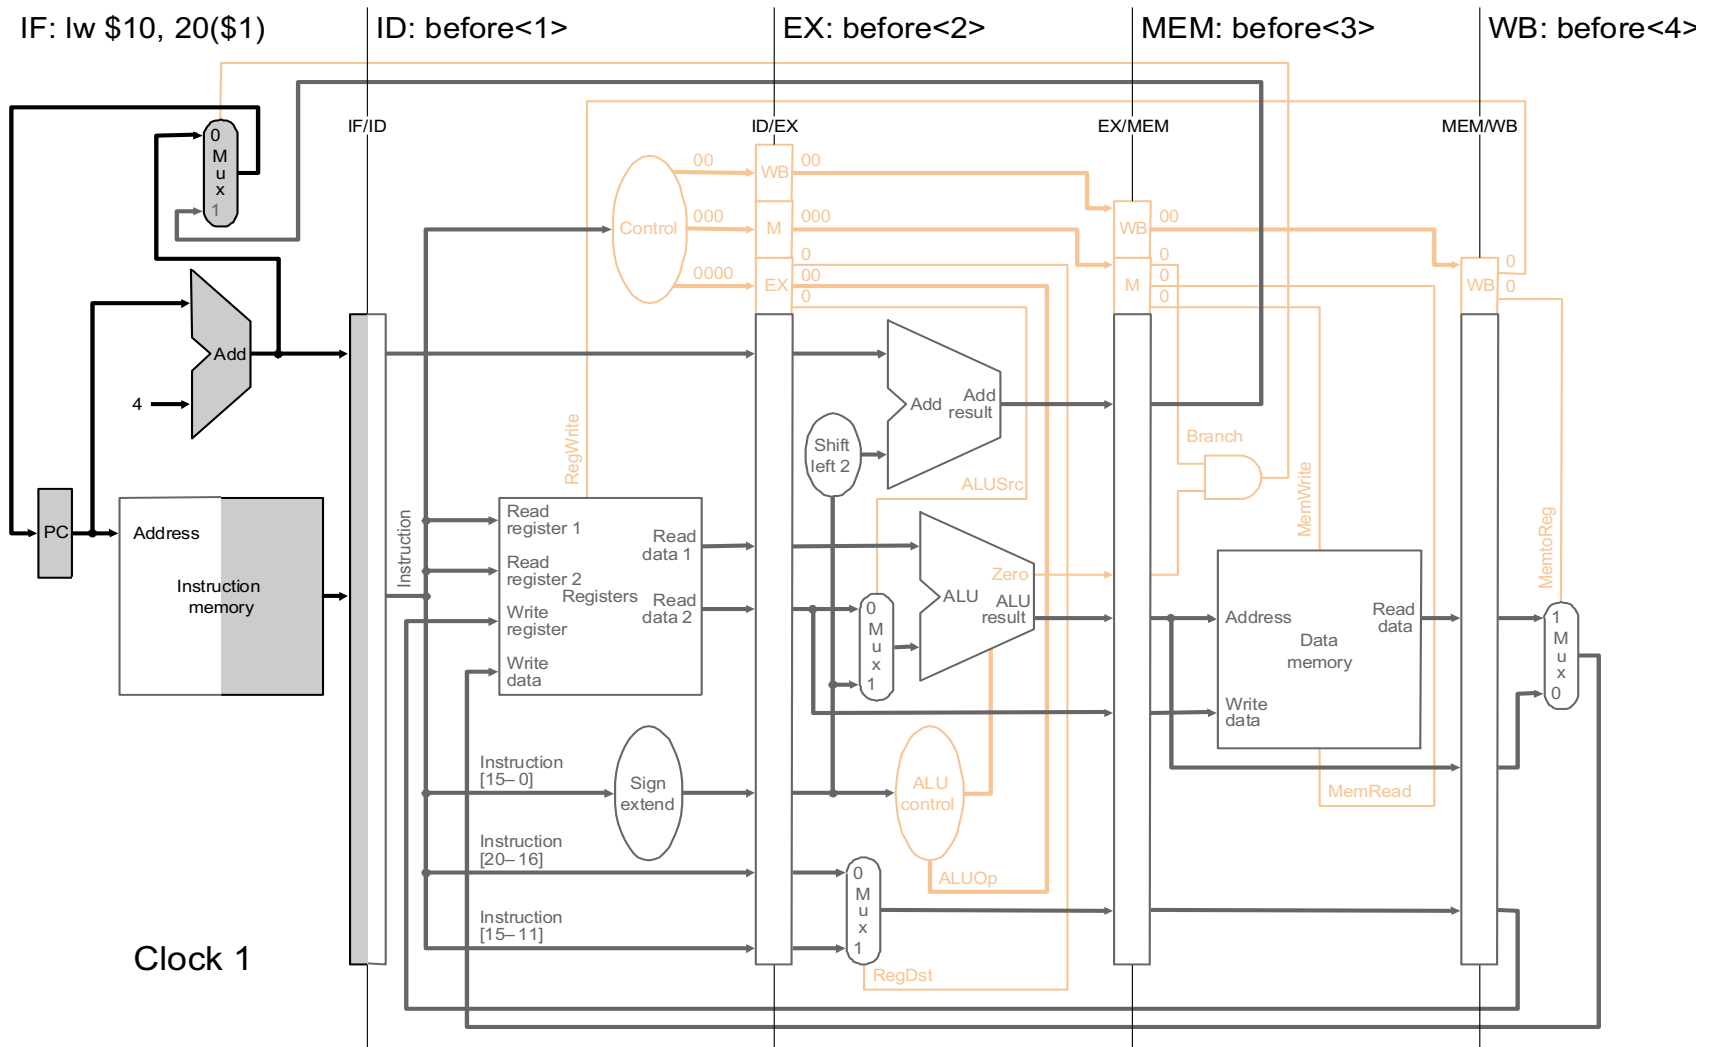
**sub**     **\$11, \$2, \$3**

**and**     **\$12, \$4, \$5**

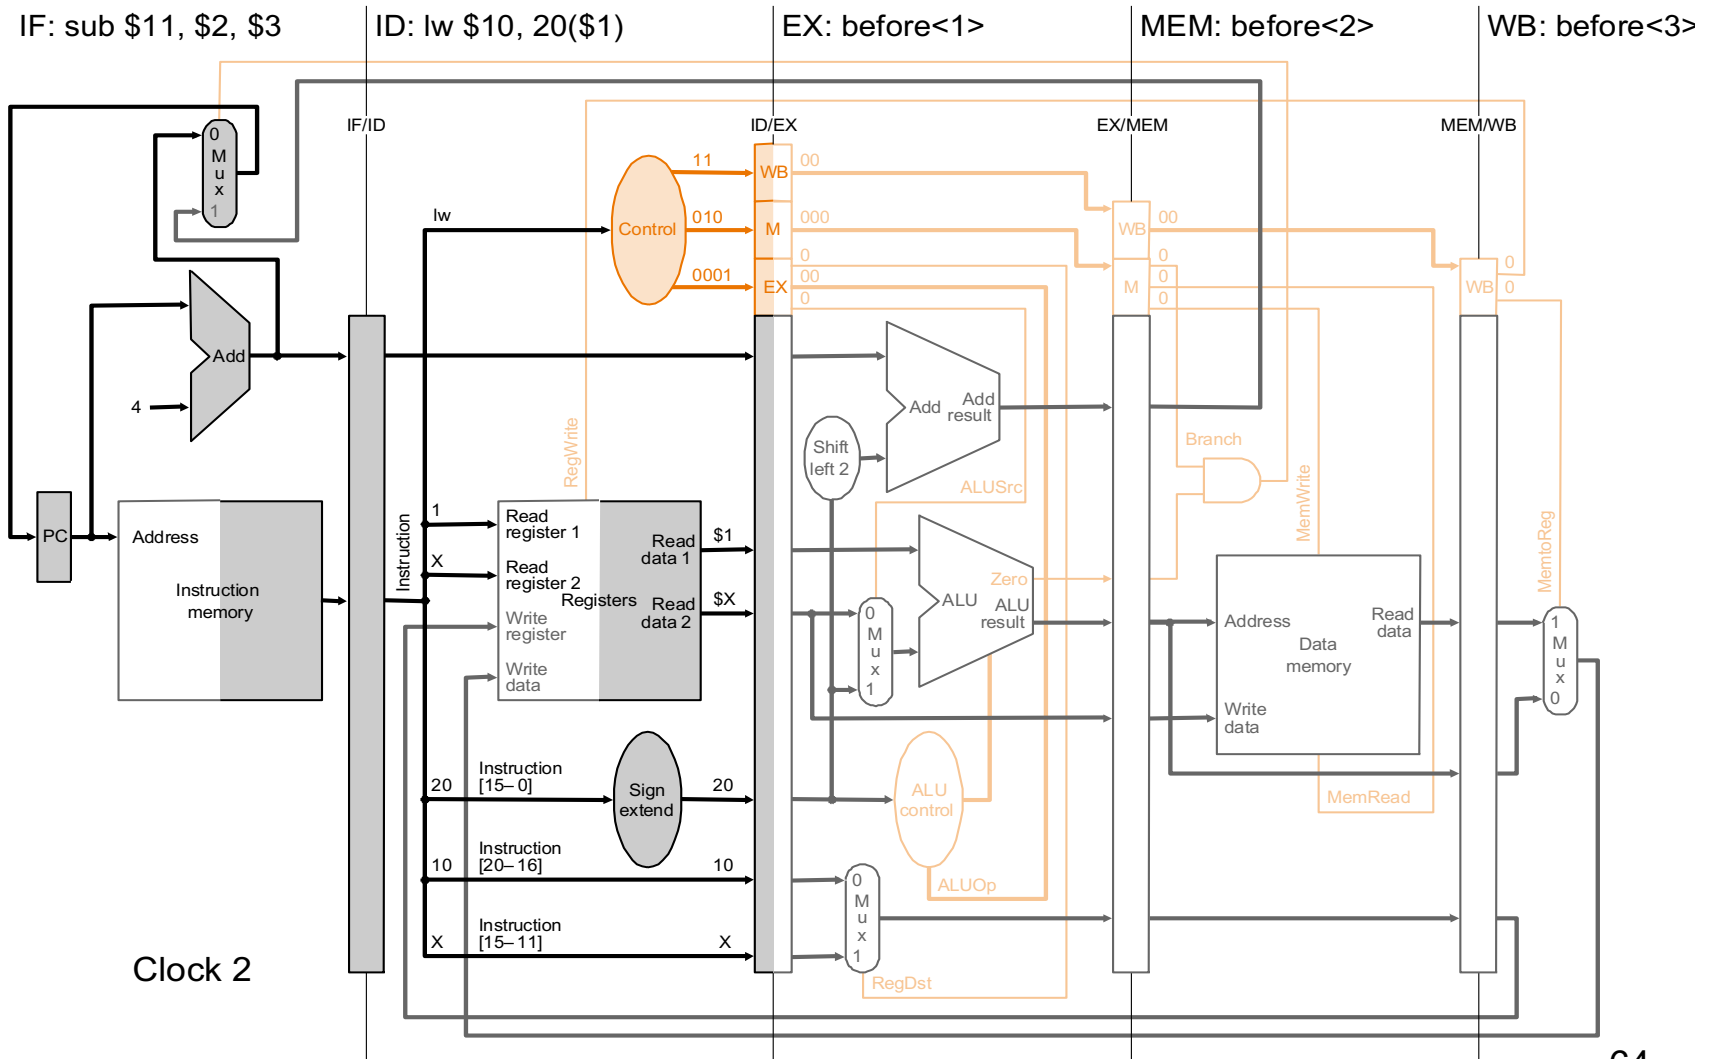
**or**      **\$13, \$6, \$7**

**add**     **\$14, \$8, \$9**

# Cycle 1




# Cycle 2






# Cycle 3






# Cycle 5






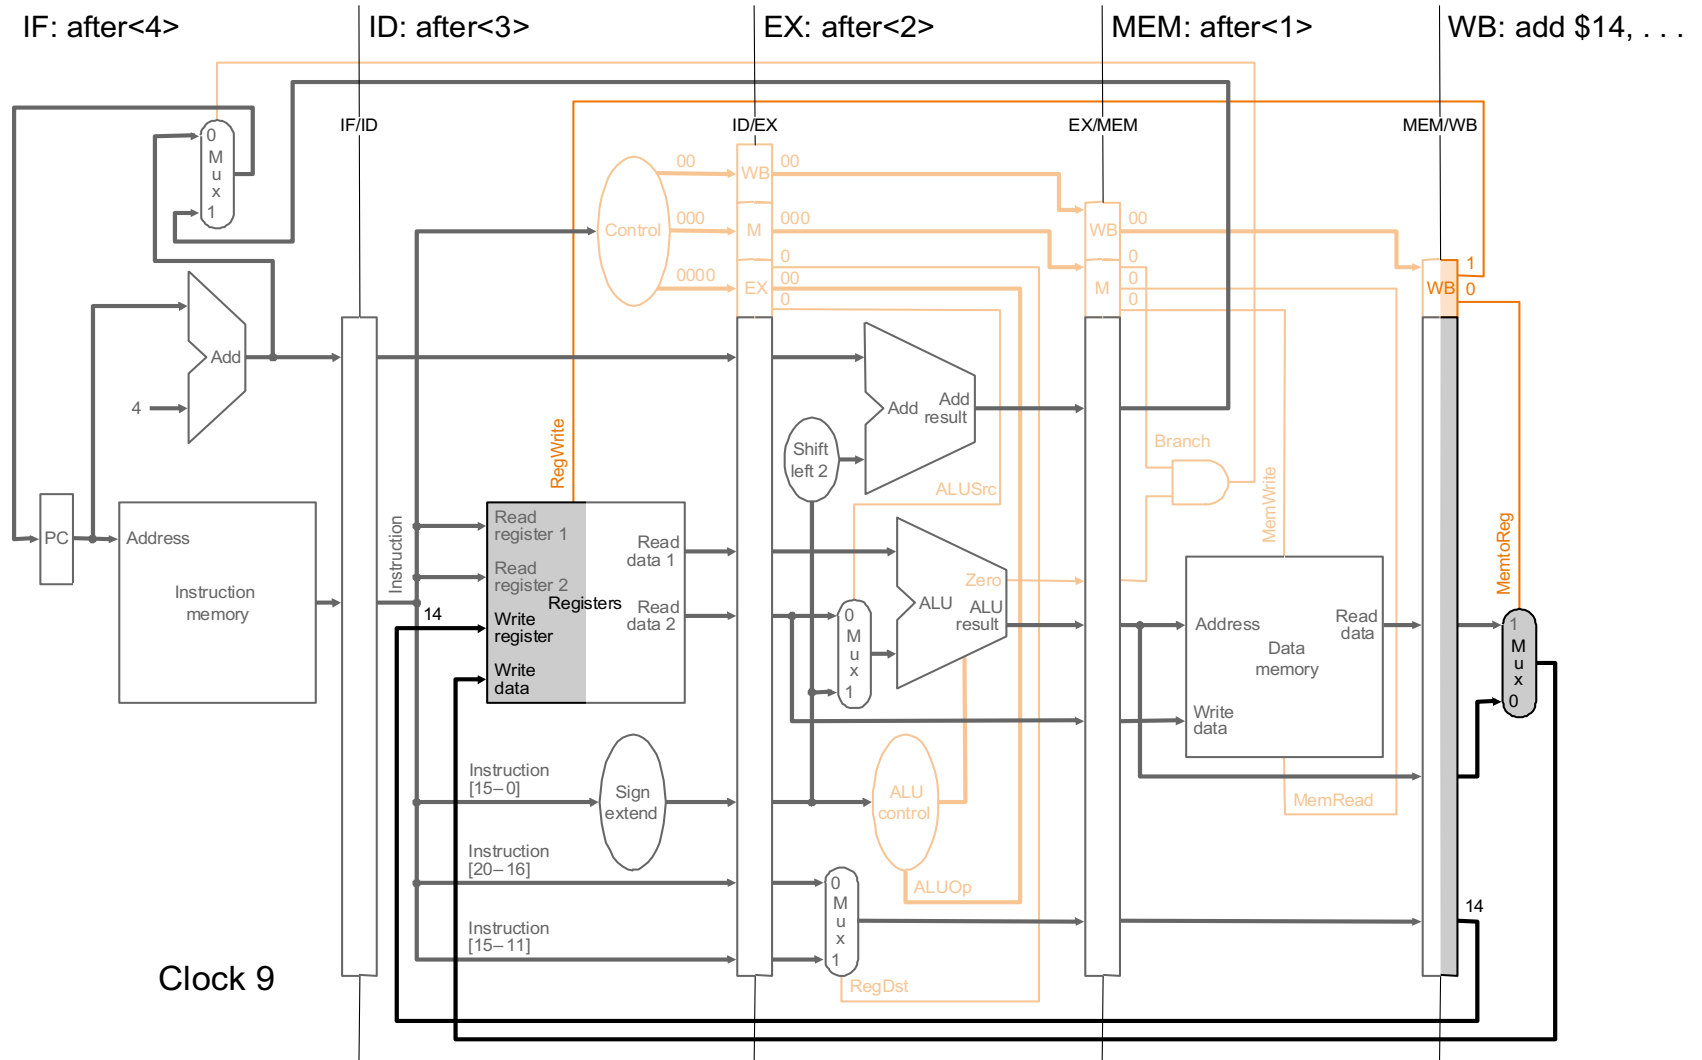
# Cycle 7



# Cycle 8



# Cycle 9



# Summary of Pipeline

---

- Pipelining is a fundamental concept
  - Multiple steps using distinct resources
  - Utilize capabilities of datapath by pipelined instruction processing
    - Start next instruction while working on the current one
    - Limited by length of longest stage (plus fill/flush)
    - Need to detect and resolve hazards
- What makes it easy in MIPS?
  - All instructions are of the same length
  - Just a few instruction formats
  - Memory operands only in loads and stores
- What makes pipelining hard? hazards